

Neural Networks

An Empirical Study on Information Extraction using Large Language Models

--Manuscript Draft--

Manuscript Number:	NEUNET-D-24-02205
Article Type:	Article
Section/Category:	Mathematical and Computational Analysis
Keywords:	Information Extraction; Large Language Models; Empirical Study
Corresponding Author:	Lu Liu Jilin University College of Software CHINA
First Author:	Ridong Han
Order of Authors:	Ridong Han
	Chaohao Yang
	Tao Peng
	Prayag Tiwari
	Xiang Wan
	Lu Liu
	Benyou Wang
Abstract:	<p>Human-like large language models (LLMs), especially the most powerful and popular ones in OpenAI's GPT family, have proven to be very helpful for many natural language processing (NLP) related tasks. Therefore, various attempts have been made to apply LLMs to information extraction (IE), which is a fundamental NLP task that involves extracting information from unstructured plain text. To demonstrate the latest representative progress in LLMs' information extraction ability, we assess the information extraction ability of GPT-4 (the latest version of GPT at the time of writing this paper) from four perspectives: Performance, Evaluation Criteria, Robustness, and Error Types. Our results suggest a visible performance gap between GPT-4 and state-of-the-art (SOTA) IE methods. To alleviate this problem, considering the LLMs' human-like characteristics, we propose and analyze the effects of a series of simple prompt-based methods, which can be generalized to other LLMs and NLP tasks. Rich experiments show our methods' effectiveness and some of their remaining issues in improving GPT-4's information extraction ability.</p>
Suggested Reviewers:	Bangzuo Zhang Northeast Normal University zhangbz@nenu.edu.cn
	Xiujuan Xu Dalian University of Technology xjxu@dlut.edu.cn
	Richong Zhang Beihang University zhangrc@act.buaa.edu.cn

Dear Editor,

We would like to submit the manuscript entitled “**An Empirical Study on Information Extraction using Large Language Models**” for publication in “**Neural Networks**”.

This paper has an original arxiv version entitled “Is Information Extraction Solved by ChatGPT? An Analysis of Performance, Evaluation Criteria, Robustness and Errors”, whose url link is “ <https://doi.org/10.48550/arXiv.2305.14450>”. To demonstrate the latest representative progress in LLMs’ information extraction ability, we assess the information extraction ability of GPT-4 (the latest version of GPT at the time of writing this paper) from four perspectives: Performance, Evaluation Criteria, Robustness, and Error Types. Our results suggest a visible performance gap between GPT-4 and state-of-the-art (SOTA) IE methods. To alleviate this problem, considering the LLMs’ human-like characteristics, we propose and analyze the effects of a series of simple prompt-based methods, which can be generalized to other LLMs and NLP tasks. Rich experiments show our methods’ effectiveness and some of their remaining issues in improving GPT-4’s information extraction ability.

The manuscript is our original work and has not been published nor has it been submitted simultaneously elsewhere. All authors have checked the manuscript and have agreed to the submission.

Thank you very much for your attention and consideration. We are looking forward to receiving your comments.

Sincerely yours,

Ridong Han, Chaohao Yang, Tao Peng, Prayag Tiwari, Xiang Wan, Lu Liu, Benyou Wang

An Empirical Study on Information Extraction using Large Language Models

Ridong Han^{a,b,d,e,1}, Chaohao Yang^{e,1}, Tao Peng^{a,b,c,*}, Prayag Tiwari^f, Xiang
Wan^{d,e}, Lu Liu^{b,c,*}, Benyou Wang^{d,e,*}

^aCollege of Computer Science and Technology, Jilin University, China

^bKey Laboratory of Symbolic Computation and Knowledge Engineering of Ministry of
Education, Jilin University, China

^cCollege of Software, Jilin University, China

^dShenzhen Research Institute of Big Data, The Chinese University of Hong Kong,
Shenzhen, China

^eSchool of Data Science, The Chinese University of Hong Kong, Shenzhen, China

^fSchool of Information Technology, Halmstad University, Sweden

* Corresponding Authors

Email addresses: tpeng@jlu.edu.cn (Tao Peng), liulu@jlu.edu.cn (Lu Liu),
wangbenyou@cuhk.edu.cn (Benyou Wang)

¹Ridong Han and Chaohao Yang contributed equally, and they are co-first authors.

This article has an original arxiv version entitled “Is Information Extraction Solved by ChatGPT? An Analysis of Performance, Evaluation Criteria, Robustness and Errors”, whose url link is “<https://doi.org/10.48550/arXiv.2305.14450>”.

An Empirical Study on Information Extraction using Large Language Models

Ridong Han^{a,b,d,e,1}, Chaohao Yang^{e,1}, Tao Peng^{a,b,c,*}, Prayag Tiwari^f, Xiang Wan^{d,e}, Lu Liu^{b,c,*}, Benyou Wang^{d,e,*}

^a*College of Computer Science and Technology, Jilin University, China*

^b*Key Laboratory of Symbolic Computation and Knowledge Engineering of Ministry of Education, Jilin University, China*

^c*College of Software, Jilin University, China*

^d*Shenzhen Research Institute of Big Data, The Chinese University of Hong Kong, Shenzhen, China*

^e*School of Data Science, The Chinese University of Hong Kong, Shenzhen, China*

^f*School of Information Technology, Halmstad University, Sweden*

Abstract

Human-like large language models (LLMs), especially the most powerful and popular ones in OpenAI's GPT family, have proven to be very helpful for many natural language processing (NLP) related tasks. Therefore, various attempts have been made to apply LLMs to information extraction (IE), which is a fundamental NLP task that involves extracting information from unstructured plain text. To demonstrate the latest representative progress in LLMs' information extraction ability, we assess the information extraction ability of GPT-4 (the latest version of GPT at the time of writing this paper) from four perspectives: Performance, Evaluation Criteria, Robustness, and Error Types. Our results suggest a visible performance gap between GPT-4 and state-of-the-art (SOTA) IE methods. To alleviate this problem, considering the LLMs' human-like characteristics, we propose and analyze the effects of a series of simple prompt-based

*Corresponding authors.

Email addresses: tpeng@jlu.edu.cn (Tao Peng), liulu@jlu.edu.cn (Lu Liu), wangbenyou@cuhk.edu.cn (Benyou Wang)

¹Ridong Han and Chaohao Yang contributed equally, and they are co-first authors.

This article has an original arxiv version entitled "Is Information Extraction Solved by ChatGPT? An Analysis of Performance, Evaluation Criteria, Robustness and Errors", whose url link is "<https://doi.org/10.48550/arXiv.2305.14450>".

methods, which can be generalized to other LLMs and NLP tasks. Rich experiments show our methods’ effectiveness and some of their remaining issues in improving GPT-4’s information extraction ability.

Keywords: Information Extraction, Large Language Models, Empirical Study

1. Introduction

The rapidly evolving field of natural language processing (NLP) witnesses the rise of large language models (LLMs), such as GPT-3[1], LaMDA[2], PaLM[3], etc., which have revolutionized various downstream tasks with in-context learning (ICL) [1] and chain-of-thought (COT) prompting [4]. Excitingly, just by providing appropriate instructions [5, 6] or chain-of-thought prompts [4], LLMs can achieve amazing performance on the zero-shot and few-shot scenarios of unseen tasks, even without updating parameters.

Currently, one of the most popular and powerful LLM series is OpenAI’s GPT series, which is best known for its two latest members, GPT-3.5 and GPT-4 [7], by which ChatGPT is powered [8]. These two models have exhibited powerful dialogue ability and stimulated the research boom for investigating the capabilities of LLMs. For example, Jiao et al. [9] evaluate the machine translation capability of ChatGPT, and Bang et al. [10] assess the reasoning capability of ChatGPT. As a fundamental natural language understanding task, Information extraction (IE) aims to identify structured information of interest from unstructured plain text. Its results directly affect the subsequent downstream tasks, such as question-answering [11, 12] and knowledge graph construction [13]. Besides, the LLMs’ ability to recognize target information can directly reflect their performance in understanding task instructions to generate responses. This paper, therefore, aims to conduct an empirical study on information extraction using LLMs and to demonstrate the latest representative progress in LLMs’ capabilities in information extraction. As the latest GPT version when writing this paper, GPT-4 has shown powerful capabilities beyond previous LLMs, including GPT-3.5. Hence, it is reasonable to select GPT-4 as a representative

26 case of LLMs for research.

27 In this paper, we evaluate GPT-4’s capabilities on IE tasks in terms of
28 four perspectives, including **Performance**, **Evaluation Criteria**, **Robust-**
29 **ness**, and **Error Types**. In response to the performance gaps revealed in the
30 evaluation, we also proposed three prompt-based performance **Improvement**
31 **Methods** considering the human-like characteristics of GPT-4.

32 ***Performance.*** We evaluate the performance of GPT-4 on 16 datasets with 14
33 IE sub-tasks under 3 settings: zero-shot prompts, few-shot ICL prompts, and
34 few-shot COT prompts. We also evaluate the performance of GPT-3.5 on these
35 settings for more comparison, which reveals the performance improvement of
36 the GPT series on IE tasks². The results indicate the following conclusions:

- 37 • There is a significant performance gap between the two GPT models and
38 SOTA methods. The harder the task, the larger the gap. However, the
39 two GPT models can equal or exceed SOTA methods on a few simple
40 tasks.
- 41 • The performance of GPT-4 is better than that of GPT-3.5 in most tasks.
- 42 • Using few-shot ICL prompts generally leads to significant improvements,
43 but still visibly lags behind SOTA results, while the chain-of-thought
44 prompting cannot guarantee further gains compared to few-shot ICL prompts.

45 ***Evaluation Criteria.*** Through the manual checking, we find that GPT-3.5
46 and GPT-4 tend to identify spans that contain or are contained by the anno-
47 tated ones, i.e., the recognized spans usually contain qualifiers such as crowns,
48 quantifiers, adjectives, time, place, etc. Thus, the previous span hard-matching
49 strategy is not suitable for the evaluation of LLMs like GPTs that generate
50 human-like responses. We propose a soft-matching strategy to solve this prob-
51 lem and display evaluation results for LLMs more accurately.

²We use gpt-3.5-turbo-0314 and gpt-4-0125-preview for research, the latter being the latest GPT version when writing this paper.

52 **Robustness.** We conduct comparisons and analysis on three dimensions: In-
53 valid Output, Frequency of Target Types, and The Order of Entities. We find
54 that:

- 55 • GPT-4 rarely outputs invalid responses in most cases.
- 56 • The frequency of target types has a significant impact on GPT-4’s perfor-
57 mance.
- 58 • GPT-4 is sensitive to the order of entities, and can understand the subject-
59 object relationships of entities better than GPT-3.5, but still needs further
60 improvement.

61 **Error Types.** We summarize 7 types of errors on IE tasks by manually check-
62 ing ChatGPT’s responses, including Missing spans, Missing types, Unmentioned
63 spans, Unannotated spans, Incorrect span offsets, Undefined types, and Incor-
64 rect types. We find that “Missing spans” and “Unannotated spans” are the
65 most dominant error types, accounting for more than 60% of errors in most
66 cases. The widespread presence of “Unannotated spans” also raises concerns
67 about the quality of the annotated data. Maybe using GPT-4 to assist in an-
68 notating data is a better solution.

69 **Improvement Methods.** As a generative large language model, GPT-4’s mode
70 in solving NLP tasks including IE is similar to that of humans. It outputs an-
71 swers token by token based on the requirements specified by the instruction and
72 input text given in the prompt, during which the rich knowledge learned from
73 the pre-training corpus will be used. Therefore, we can consider unfine-tuned
74 GPT-4 as a knowledgeable information extraction layman, whose IE ability can
75 be made more proficient by designing prompts in the same way as to inspire
76 human laymen, and this method should also apply to fine-tuned models. Based
77 on the above analysis, we propose the following three prompt design methods:

- 78 • **Task-related Knowledge Informing:** Inform the model of the task-
79 related knowledge required to perform the task, including The meanings
80 of task-related terms such as “trigger” and “argument” in IE tasks.

- **Methodology Specifying:** Give a specific operation methodology to make the model more proficient in information extraction.
- **Sufficient Extraction Reminder:** When the amount of information extracted is insufficient, explicitly remind the model to extract all the required information in the text it can.

The effectiveness of these methods has been verified via rich experiments. Since the first method is the most basic one for performing tasks, it is directly applied to the performance evaluation part, and we will evaluate the performance without this method in this part for comparison.

2. Related Work

Large Language Models. Based on the highly parallelizable Transformer architecture [14], pre-trained language models (PLMs) such as BERT [15], BART [16], etc., have shown powerful capabilities to solve a wide variety of NLP tasks. Some researchers find that scaling PLMs by increasing model size or data size often leads to more powerful capabilities, as long as the scaling law is followed [17, 18]. Thus, numerous large-size models have been proposed, such as GPT-3 [1], LaMDA [2], MT-NLG [19], PaLM [3] and GPT-4 [7], which typically have more than 100 billion parameters. The NLP community refers to these large-size PLMs as large language models (LLMs). Unlike small-sized PLMs, LLMs usually exhibit amazing emergent abilities [20, 21] that enable them to achieve good performance in zero-shot and few-shot scenarios of unseen tasks, as long as the appropriate instructions [22, 23, 24] or chain-of-thought prompts [4] are provided.

GPT Series. One of the best-known examples of LLMs is OpenAI’s GPT (Generative Pre-Training Transformer) series, including GPT-1 [25], GPT-2 [26], GPT-3 [1], GPT-4 [7], etc. A key milestone in the development process is InstructGPT [6], a framework for instruction fine-tuning based on reinforcement

learning from human feedback (RLHF) [27]. The framework allows a large language model to be adapted to a large number of NLP tasks simultaneously, and leverages human feedbacks to align the model output with human preferences in order to generate responses more consistent with human expectations. As the successor of InstructGPT, ChatGPT, powered by GPT-3.5 and GPT-4, has exploded the field of artificial intelligence (AI), and attracted an unprecedented wave of enthusiasm. It can interact with humans through multiple turns of dialogue, understand user intent, accomplish instructions, and return human-like responses. Shocked by ChatGPT’s performance, some papers already consider GPT-4 as an early version of artificial general intelligence (AGI) [28, 29].

Information Extraction. As a popular and fundamental task, information extraction (IE) aims to extract structured knowledge of interest from unstructured plain text. The output mainly includes entities, relations between entities, event arguments, opinions, human sentiments, etc. Due to the different target information, IE mainly involves 4 tasks, including named entity recognition (NER) [30, 31, 32, 33], relation extraction (RE) [34, 35, 36, 37, 38, 39, 40], event extraction (EE) [41, 42, 43] and aspect-based sentiment analysis (ABSA) [44, 45, 46, 47, 48, 49]. Since the result of IE directly affects the performance of subsequent higher-level applications, the importance of IE cannot be overstated. This paper intends to evaluate the performance of GPT-4 on IE, in detail.

Evaluation of ChatGPT. ChatGPT’s powerful dialog capability has triggered widespread research interest in the fields of NLP and LLM. Since ChatGPT is based on two closed models and no training details are provided, researchers are exploring its concerns and capabilities. The concerns involve ethical risks [50, 51], patient privacy [52], fabricated misinformation [53, 54], education integrity [55] and legal challenges [56]. For its capabilities, researchers evaluate the performance of ChatGPT on different tasks, including stance detection [57], question-answering [58], machine translation [9], sentiment analysis [59] and other general NLP tasks [60, 61, 62, 10]. In addition, for the information extraction task, Wei et al. [63] propose a two-stage framework, ChatIE, to use

138 ChatGPT for zero-shot information extraction, and evaluate its performance in
 139 detail. Li et al. [64] measure the performance, explainability, calibration, and
 140 faithfulness of ChatGPT on IE tasks. Following these works, this paper focuses
 141 on demonstrating the latest progress in LLMs’ capabilities in information ex-
 142 traction represented by ChatGPT-related models. we measure the performance
 143 of the latest version of GPT-4 on multiple datasets of 14 IE subtasks, explore
 144 the impact of in-context learning (ICL) and chain-of-thought (COT) prompts
 145 on performance, evaluate robustness by scenario, and analyze error types. Fol-
 146 lowing Wei et al. [63], we also propose three novel improvement methods for
 147 GPT-4’s IE performance based on its human-like characteristics. Our perspec-
 148 tive differs significantly from Li et al. [64], and we evaluate more IE sub-tasks
 149 on more benchmarks.

150 **3. Experimental Protocol**

151 *3.1. Tasks*

152 In this paper, we consider 4 well-representative IE tasks, including Named
 153 Entity Recognition (NER), Relation Extraction (RE), Event Extraction (EE),
 154 and Aspect-based Sentiment Analysis (ABSA). Since each task contains several
 155 sub-tasks or scenarios, we conduct evaluations and analysis on the following 14
 156 sub-tasks:

- 157 • Flat Entity Recognition (NER-Flat): Recognizing all entities within the
 158 text. Each entity is identified as a separate entity, without any hierarchical
 159 relationship between them.
- 160 • Nested Entity Recognition (NER-Nested): Recognizing all entities within
 161 the text. Each entity can be nested inside other entities, i.e., an entity
 162 may contain other sub-entities.
- 163 • Relation Classification (RE-RC): Determining the relationship between a
 164 given subject-object pair of entities within the text.

- 165 • Relational Triplet Extraction (RE-Triplet): Identifying entities and their
166 relationships simultaneously in the form of (subject entity, relation, object
167 entity) triplets.
- 168 • Event Detection (EE-Trigger): Identifying the word or phrase that indi-
169 cates the occurrence of an event, and categorizing its corresponding event
170 type.
- 171 • Event Argument Extraction (EE-Argument): Recognizing the entities in-
172 volved in the given event, and classifying their corresponding roles.
- 173 • Trigger-Argument joint Extraction (EE-Joint): Identifying event trigger,
174 event type, and all arguments with their roles simultaneously.
- 175 • Aspect Extraction (ABSA-AE): Extracting all the aspect terms from a
176 review.
- 177 • Opinion Extraction (ABSA-OE): Extracting all the opinion terms from a
178 review.
- 179 • Aspect-level Sentiment Classification (ABSA-ALSC): Predicting the sen-
180 timent polarities for every given aspect term in a review.
- 181 • Aspect-oriented Opinion Extraction (ABSA-AOE): Extracting the paired
182 opinion terms for every given aspect term in a review.
- 183 • Aspect Extraction and Sentiment Classification (ABSA-AESC): Extract-
184 ing the aspect terms as well as the corresponding sentiment polarities
185 simultaneously.
- 186 • Pair Extraction (ABSA-Pair): Extracting the aspect terms as well as the
187 corresponding opinion terms simultaneously.
- 188 • Triplet Extraction (ABSA-Triplet): Extracting all aspects terms with
189 their corresponding opinion terms and sentiment polarity simultaneously.

190 3.2. Datasets

191 We select at least three datasets for each IE task, with a total of 16 datasets
192 as follows:

- 193 • For NER task, the datasets include CoNLL03 [65], FewNERD[32], ACE04[66],
194 ACE05-Ent[67], and GENIA[68].
- 195 • For RE task, the datasets include CoNLL04 [69], NYT-multi [70], and
196 SemEval 2010 [71].
- 197 • For EE task, the datasets include ACE05-Evt [67], ACE05+ [41], CASIE
198 [72], and Commodity News EE [73].
- 199 • For ABSA task, the datasets include D_{17} [74], D_{19} [75], D_{20a} [76], and D_{20b}
200 [77], which are all originated from the SemEval Challenges [78, 79, 80].

201 3.3. Prompts

202 The prompts designed in this paper all consist of five main elements: task
203 instruction, candidate target labels, output format description, demonstration
204 examples, and input text. The task instruction describes the specific IE sub-
205 task, candidate target labels are the types of target information, such as entity
206 types, relation types, etc. The output format description specifies the format of
207 outputs to facilitate the distinguishing of target information. The demonstration
208 examples exist under the few-shot In-context Learning setting, which can also
209 provide the chain-of-thought explanation. The input text is a text or review
210 from which target information is to be extracted. An example of prompts for
211 the NER-Flat sub-task is shown in Figure 1.

212 For the demonstration examples, we randomly select them from the training
213 set of each dataset in Section 3.2. To obtain the chain-of-thought prompts, we
214 construct them manually with the help of ChatGPT to generate explanations.

An example of NER-Flat sub-task prompts

Considering 4 types of named entities including “organization”, “person”, “location” and “miscellaneous”, recognize all named entities in the given sentence.

Answer in the format of '[“entity_type 1”, “entity_name 1”], [“entity_type 2”, “entity_name 2”], ...' without any explanation. If no entity exists, then just answer “[]”.

Sentence:

“Rapid Wien 5 0 5 0 3 3 5”

Answer:

[“organization”, “Rapid Wien”]

... (More examples are omitted here.)

Sentence:

“Results of semifinals in the Mahindra International squash tournament on Friday :”

Answer:

Expected Output:

[“miscellaneous”, “Mahindra International”]

Figure 1: An example of prompts for NER-Flat sub-task on CoNLL03 dataset. See the Appendix B for more prompts.

215 3.4. Setup

216 To conduct a thorough evaluation of GPT-4’s capabilities, for each IE sub-
217 task, we first measure the performance of the zero-shot scenario. Then, we
218 investigate the impact of few-shot in-context learning (ICL) and few-shot chain-
219 of-thought (COT) prompting on the performance. For the construction of few-
220 shot ICL prompts, we use the zero-shot prompt as the basic component and add
221 randomly selected samples from the corresponding training set. For few-shot
222 COT prompts, we add the chain-of-thought explanations to the few-shot ICL
223 prompts, where the chain-of-thought explanations are manually constructed via
224 dialog with ChatGPT. To eliminate the randomness of selected samples, we

select three different groups of samples for construction and report their mean performance.

We use the official API to generate all outputs from GPT-4. To prevent the influence of dialogue history, we generate the response separately for each testing sample. Unlike other work where only 30-50 samples are selected for evaluation[9, 63], we use the entire deduplicated test set of most datasets in Section 3.2 for evaluation. Too few samples will lead to low coverage and high randomness of results, too many samples are limited by the rate and expense of accessing OpenAI’s API. Since most of the datasets we use have a test set with less than 2000 samples, we limit the number of samples to a maximum of 2000 through random sampling.

Besides, we compare GPT-4, GPT-3.5, and the state-of-the-art result for each sub-task. Since calling the API of GPT-3.5 costs much less than that of GPT-4, we set its test set length limit to 3000 and select five groups of samples for few-shot ICL and COT. For the metric, we use Micro-F1 for all sub-tasks. The detailed specifications for each sub-task are as follows:

- **ABSA-AE, ASBA-OE**: An aspect/opinion is correct if its span matches the reference aspect/opinion mention.
- **ABSA-ALSC**: A sentiment polarity is correct if it matches the reference polarity of the given aspect term.
- **ABSA-AOE**: An opinion is correct if its span matches the reference opinion of the given aspect term.
- **ABSA-AESC**: An aspect-sentiment pair is correct if its aspect span and corresponding sentiment polarity are all correct.
- **ABSA-Pair**: An aspect-opinion pair is correct if its aspect span and opinion span all match the reference pair.
- **ABSA-Triplet**: A triplet is correct if its aspect span, opinion span, and corresponding sentiment polarity are all correct.

- 253 • **Flat-NER, Nested-NER:** A predicted entity is correct if its offsets and
254 type match a reference entity.
- 255 • **RE-RC:** A predicted relation is correct if its relation type matches the
256 reference type.
- 257 • **RE-Triplet:** A predicted relational triplet is correct if its relation type is
258 correct and the subject and object entity spans are all correct. We only
259 report the F1 value of relational triplets.
- 260 • **EE-Trigger:** A predicted event trigger is correct if its span and event
261 type all match the reference trigger.
- 262 • **EE-Argument:** For a given event type, a predicted argument is correct
263 if its span and role type all match the reference argument mention of this
264 event.
- 265 • **EE-Joint:** A predicted argument is correct if its span, role type, and
266 event type all match the reference argument mention. We only report the
267 F1 value of the (event type, event argument, role) triplets contained in
268 the output.

269 4. The Performance

270 In this section, we report the performance of GPT-4 and GPT-3.5 on 14
271 different sub-tasks, as shown in Table 1.

272 4.1. Performance Gap in Zero-shot Scenario

273 From the zero-shot result of Table 1, we can draw the following conclusions:

274 (1) **There is a significant performance gap between the two GPT**
275 **models and SOTA methods.** This seems obvious and reasonable since all
276 SOTA methods are trained on corresponding datasets. In other words, they are
277 fully supervised models and are not zero/few-shot ones.

278 (2) **The harder the task, the larger the performance gap.** From the
279 perspective of the four IE tasks of NER, RE, EE, and ABSA, it can be seen

Task	Dataset	SOTA	Zero-shot		ICL	COT		Ratio@SOTA
			GPT-4 (3.5)	Δ F1 (%)	GPT-4 (3.5)	GPT-4 (3.5)		
ABSA-AE	$D_{17-14lap}$	85.3[81]	46.06 (43.03)	7.04%	52.85 (48.19)	55.92 (54.50)	54.0%	
	$D_{17-14res}$	87.4[81]	64.74 (55.65)	16.34%	73.84 (70.99)	75.71 (72.41)	74.1%	
	$D_{17-15res}$	79.4[81]	48.71 (40.33)	20.78%	57.81 (53.49)	58.37 (59.27)	61.3%	
ABSA-OE	$D_{17-14lap}$	84.4[81]	51.44 (48.45)	6.17%	57.34 (57.89)	39.26 (50.78)	60.9%	
	$D_{17-14res}$	88.9[81]	62.68 (59.48)	5.38%	70.24 (71.61)	61.92 (58.74)	70.5%	
	$D_{17-15res}$	82.7[81]	49.49 (46.39)	6.69%	57.04 (53.96)	50.77 (47.11)	59.8%	
ABSA-ALSC	$D_{17-14lap}$	76.8[45]	79.82 (74.56)	7.05%	80.58 (76.76)	79.92 (75.35)	103.9%	
	$D_{17-14res}$	82.0[82]	86.46 (81.16)	6.53%	85.16 (81.85)	87.37 (79.77)	105.4%	
	$D_{17-15res}$	74.9[44]	87.06 (88.13)	-1.21%	87.49 (86.47)	88.60 (75.23)	116.2%	
ABSA-AOE	$D_{19-14lap}$	82.2[46]	56.57 (57.60)	-1.78%	62.17 (64.83)	56.14 (55.43)	68.8%	
	$D_{19-14res}$	86.4[46]	67.02 (67.67)	-0.96%	71.69 (71.60)	62.45 (66.42)	77.6%	
	$D_{19-15res}$	81.6[46]	60.57 (67.03)	-9.64%	64.83 (70.29)	55.74 (59.60)	74.2%	
	$D_{19-16res}$	89.2[46]	65.13 (73.27)	-11.10%	71.64 (78.23)	60.09 (68.44)	73.0%	
ABSA-AESC	$D_{20a-14lap}$	70.1[49]	44.77 (45.48)	-1.56%	52.70 (49.50)	50.88 (48.87)	63.9%	
	$D_{20a-14res}$	79.7[49]	59.94 (59.08)	1.45%	65.93 (65.98)	63.69 (64.82)	75.2%	
	$D_{20a-15res}$	71.6[49]	61.33 (53.91)	13.77%	65.71 (63.66)	64.74 (66.07)	85.7%	
	$D_{20a-16res}$	77.5[49]	56.51 (55.40)	2.01%	63.38 (63.11)	65.88 (65.93)	72.9%	
ABSA-Pair	$D_{20a-14lap}$	69.1[81]	38.20 (31.76)	20.29%	42.80 (41.59)	40.91 (35.75)	55.3%	
	$D_{20a-14res}$	77.8[83]	48.15 (50.05)	-3.79%	59.70 (58.88)	55.42 (49.82)	61.9%	
	$D_{20a-15res}$	69.4[49]	50.05 (44.41)	12.69%	55.52 (53.76)	51.77 (49.62)	72.1%	
	$D_{20a-16res}$	78.2[81]	49.91 (50.20)	-0.58%	57.93 (58.88)	57.52 (51.64)	63.8%	
ABSA-Triplet	$D_{20b-14lap}$	61.7[48]	35.49 (33.17)	6.99%	40.18 (39.01)	36.30 (33.18)	57.5%	
	$D_{20b-14res}$	74.4[48]	53.20 (41.50)	28.20%	56.66 (54.89)	52.66 (48.90)	71.5%	
	$D_{20b-15res}$	66.1[48]	51.26 (38.89)	31.82%	51.31 (47.88)	46.63 (46.55)	77.6%	
	$D_{20b-16res}$	72.3[48]	54.82 (47.67)	15.00%	57.70 (56.55)	54.49 (51.84)	75.8%	
NER-Flat	CoNLL03	94.6[31]	72.30 (60.10)	20.30%	78.50 (70.53)	76.19 (74.63)	76.4%	
	FewNERD	67.1[32]	47.84 (31.56)	51.59%	49.15 (36.87)	50.61 (46.55)	71.3%	
NER-Nested	ACE04	88.5[33]	31.43 (27.80)	13.05%	43.35 (38.52)	46.03 (40.57)	35.5%	
	ACE05-Ent	87.5[33]	24.68 (23.38)	5.54%	44.68 (36.17)	41.46 (33.98)	28.2%	
	GENIA	81.5[33]	46.22 (38.09)	21.36%	56.97 (48.82)	54.57 (50.89)	56.7%	
RE-RC	CoNLL04	-	82.07 (59.21)	38.61%	92.92 (55.32)	-	-	
	NYT-multi	93.5[37]	47.79 (30.96)	54.36%	54.47 (26.88)	-	51.1%	
	SemEval2010	91.3[35]	44.80 (39.27)	14.09%	44.98 (39.44)	-	49.1%	
RE-Triplet	CoNLL04	78.8[84]	26.53 (17.84)	48.68%	37.28 (24.30)	35.88 (11.09)	33.7%	
	NYT-multi	86.8[40]	12.74 (3.48)	266.23%	19.55 (12.24)	13.41 (2.33)	14.7%	
	SemEval2010	73.2[85]	5.09 (5.82)	-12.55%	16.80 (12.85)	-	7.0%	
EE-Trigger	ACE05-Evt	77.1[85]	41.62 (17.55)	137.18%	47.30 (27.33)	27.40 (7.81)	54.0%	
	ACE05+	72.8[41]	43.24 (18.22)	137.34%	47.17 (29.17)	26.59 (9.06)	59.4%	
	CASIE	72.0[86]	24.98 (7.24)	245.03%	27.17 (18.23)	16.92 (3.95)	34.7%	
	Commodity News EE	94.0[42]	47.61 (17.90)	165.98%	54.53 (37.79)	27.63 (12.75)	50.6%	
EE-Argument	ACE05-Evt	73.5[43]	29.43 (25.09)	17.31%	34.40 (31.62)	-	40.0%	
	ACE05+	73.0[43]	29.72 (25.80)	15.19%	36.32 (32.02)	-	40.7%	
	CASIE	-	30.99 (17.31)	79.01%	31.32 (27.35)	-	-	
	Commodity News EE	-	16.57 (12.06)	37.38%	24 (15.08)	-	-	
EE-Joint	ACE05-Evt	57.3[86]	3.70 (8.74)	-57.62%	25.52 (13.82)	-	6.5%	
	ACE05+	56.8[43]	3.51 (10.12)	-65.28%	26.19 (13.33)	-	6.2%	
	CASIE	63.5[85]	16.78 (14.24)	17.81%	27.75 (18.96)	-	26.4%	
	Commodity News EE	90.0[42]	2.14 (8.46)	-74.67%	20.77 (14.02)	-	2.4%	

Table 1: The performances of GPT-4 and GPT-3.5 on different datasets over multiple standard IE tasks. “ Δ F1 (%)” indicates the zero-shot performance improvement rate of GPT-4 compared to GPT-3.5, and “Ratio@SOTA” indicates the percentage value of GPT-4 performance vs. SOTA in the zero-shot scenario.

that ABSA tasks perform significantly better than RE and EE tasks. Almost all sub-tasks of ABSA for GPT-3.5 can reach more than 50% of SOTA, while all sub-tasks of RE and EE rarely exceed 30% of SOTA. For GPT-4, these two percentages are respectively 60% and 50%. One reason is that ABSA tasks involve only aspect and opinion terms and are much simpler, While RE and EE tasks involve many target types and are much harder. For example, there are 24 relation types in the NYT-multi dataset.

(3) **The harder the scenario, the larger the performance gap.** Each IE task has several scenarios. For NER tasks, the NER-Flat scenario is intuitively simpler than NER-Nested, and the performance of NER-Flat is significantly better than NER-Nested. For other tasks, including RE, EE, and ABSA, we can observe similar results.

(4) **On a few simple cases, the two GPT models can equal or exceed the performance of SOTA methods.** We can find that both models can achieve comparable performance with SOTA methods on the ABSA-ALSC sub-task, and can even surpass SOTA results, reaching up to 117.7% (GPT-3.5) and 116.2% (GPT-4) of SOTA. The sub-task is a simple sentiment classification, and the candidate polarities include “positive”, “neutral” and “negative”.

(5) **The performance of GPT-4 is better than GPT-3.5 in most tasks.** For 36 out of all 48 cases, the zero-shot scores of GPT-4 are higher than that of GPT-3.5. For the EE-Trigger sub-task, the improvement rates are even higher than 100%. Results prove the surprising progress achieved by the GPT series in recent times.

4.2. Mitigate the Gap

The observed performance gap in the above subsection is not consistent with our actual experience with ChatGPT, no matter powered by GPT-3.5 or GPT-4. To mitigate the gap, we add a few randomly selected demonstration examples to construct few-shot ICL prompts and few-shot COT prompts. We report the means of the selected example groups in Table 1.

For the few-shot ICL setting, it can be seen that “using few-shot ICL prompts

generally leads to significant improvements (about 3.0~13.0 F1 value), but still obviously lags behind SOTA results”. This seems to be inconsistent with the conclusion of Wadhwa et al. [87] that ChatGPT can achieve performance equivalent to SOTA methods by providing some demonstration examples. One reason may be that Wadhwa et al. [87] provide more demonstration examples, i.e., almost 20 examples, while we only provide 5 demonstration examples. So, with a smaller number of demonstration examples, the few-shot ICL prompts cannot radically eliminate the performance gap.

For the few-shot COT setting, we can find that “the use of few-shot COT prompts cannot guarantee further gains compared to few-shot ICL prompts, sometimes it is worse than the performance of few-shot ICL prompts”. The possible reasons are that the quality of constructed chain-of-thought prompts is not good enough and GPT-3.5 and GPT-4 are too sensitive for the few-shot COT prompts.

To sum up, we conclude that **both GPT-3.5 and GPT-4 struggle to achieve comparable performance to the corresponding SOTA methods in both zero-shot and few-shot scenarios, even if the chain-of-thought explanations are provided.**

5. Rethink the Gap

In this section, we rethink the performance gap from the perspective of evaluation criteria. Following the evaluation method of previous work [88, 84, 86], we strictly match the start and end indices of the predicted target text span (e.g., entity spans, opinion spans). This method may not be suitable for the evaluation of LLMs like GPTs that generate human-like responses.

We manually check ChatGPT’s responses, and find that **GPT-3.5 and GPT-4 tend to identify spans that contain annotated spans or are contained by annotated spans, to get closer to humans.** All sub-tasks in Section 3.1 involve four types of span: entities, event triggers, aspect terms, and opinion terms. For each type of span, we select several typical annotated

Type	Annotated Spans	Predicted Spans
Entity	PGA Europro Tour	2021 PGA Europro Tour
	University of Michigan	The University of Michigan
	Australia	Western Australia
Event Trigger	war	move toward war
	fighting	commit fighting forces
	killed	marines killed
Aspect Term	USB ports	multiple USB ports
	application	application crash
	cable	extender cable
Opinion Term	fast	super fast
	well worth	well worth it
	not handle	does not handle

Table 2: The selected annotated spans and their corresponding predicted spans.

spans and their corresponding predicted spans and show them in Table 2. It can be seen that the annotated spans usually do not contain qualifiers such as quantifiers, articles, adjectives, time, place, etc, while the spans predicted usually contain these qualifier parts, which are also correct target information. For example, “*University of Michigan*” and “*The University of Michigan*” indicate the same target information, although the offsets are different.

Therefore, to incorporate this case, we propose a soft-matching approach to obtain more accurate evaluation results, as shown in Algorithm 1, where *GetSimilarity*(\cdot) indicates a method to calculate the similarity, for which we use the *SequenceMatcher*.*ratio*(\cdot) method in the python package difflib. The similarity calculated takes the value from 0 to 1. For the threshold γ , we set it to 0.5 by default, since this process can be seen as the binary classification problem. We assume that *when the predicted span and the annotated span are only different in offset (ensured by Line 6 in the algorithm), the predicted span is reasonable and meaningful if the similarity value is higher than 0.5*.

Algorithm 1 Soft-Matching Strategy

Input: the sentence s , the list of annotated spans L_A in sentence s , a predicted span p , the similarity threshold γ .

Output: Return *True* if one of the two spans (p and the annotated span t with the highest similarity with p) contains the other and the similarity is greater than γ , otherwise return *False*.

Begin:

0. $Similarity \leftarrow []$
1. **for** t **in** L_A :
2. $score \leftarrow GetSimilarity(t, p)$
3. $Similarity.append(score)$
4. $score, max_index \leftarrow max(Similarity)$
5. $t \leftarrow L_A[max_index]$
6. **if** p contains t **or** t contains p :
7. **if** $score > \gamma$:
8. **return** *True*.
9. **return** *False*.

End.

354 We compare the evaluation results between the default hard-matching strat-
355 egy and the soft-matching strategy for related sub-tasks and show them in
356 Table 3. For space reasons, we only report the results of one dataset for each
357 sub-tasks. From Table 3, it can be seen that **the soft-matching strategy de-**
358 **livers consistent and significant performance gains, with up to 16.50**
359 **F1 value.** Interestingly, the improvement on simple sub-tasks is much more
360 noticeable, i.e., ABSA tasks have generally higher overall performance gains
361 than EE tasks. Further, although the soft-matching strategy brings significant
362 gains, it does not reach a comparable level with SOTA methods. This is still
363 consistent with the conclusions of Section 4.

Task	Dataset	SOTA	Hard	Soft	Δ F1 (%)
ABSA-AE	D_{17} -14lap	85.3	46.06	52.36	+6.30 (13.69%)
ABSA-OE	D_{17} -14lap	84.4	51.44	65.35	+13.91 (27.03%)
ABSA-AOE	D_{19} -14lap	82.2	56.57	73.07	+16.50 (29.17%)
ABSA-AESC	D_{20a} -14lap	70.1	44.77	52.55	+7.79 (17.39%)
ABSA-Pair	D_{20a} -14lap	69.1	38.20	53.7	+15.49 (40.55%)
ABSA-Triplet	D_{20b} -14lap	61.7	35.49	47.60	+12.11 (34.11%)
NER-Flat	CoNLL03	94.6	72.30	74.29	+1.99 (2.75%)
NER-Nested	ACE05-Ent	87.5	24.68	33.51	+8.83 (35.80%)
RE-Triplet	CoNLL04	78.8	26.53	34.75	+8.22 (31.00%)
EE-Trigger	ACE05-Evt	77.1	41.62	45.94	+4.31 (10.37%)
EE-Argument	ACE05-Evt	73.5	29.43	41.20	+11.77 (40.00%)
EE-Joint	ACE05-Evt	57.3	3.70	6.17	+2.47 (66.67%)

Table 3: Comparison of results between default hard-matching strategy (**Hard**) and soft-matching strategy (**Soft**). Δ indicates the performance change caused by the soft-matching strategy. ABSA-ALSC and RE-RC sub-tasks are not included in this table because their required output contains no span.

6. Robustness Analysis

6.1. Invalid Output

Since GPT-4 is a generative model, the output responses may be irrelevant information that does not meet the task requirements. In this subsection, we investigate how many invalid responses GPT-4 returns for different IE tasks. Here invalid responses refer to responses with incorrect format or unexpected content not generated as required by task-specific prompts, which make it impossible to recognize parts of or all of the response content. The frequency and ratio of recognizable error types are analyzed in Section 7. For each sub-task, we report the ratio of invalid responses under the zero-shot setting in Table 4. For convenience, we select one dataset for each sub-task as Table 3. From the results, it can be that **in most cases, GPT-4 rarely outputs invalid re-**

376 **sponses.** However, on the EE-Joint sub-task, invalid responses account for up
 377 to 11.97%, which may result from its sophisticated format description and the
 378 difficulty of the sub-task itself.

Task	Dataset	#Prompt	#Invalid	Ratio (%)
ABSA-AE	$D_{17-14lap}$	800	0	0.00%
ABSA-OE	$D_{17-14lap}$	800	0	0.00%
ABSA-ALSC	$D_{17-14lap}$	654	0	0.00%
ABSA-AOE	$D_{19-14lap}$	482	3	0.62%
ABSA-AESC	$D_{20a-14lap}$	339	0	0.00%
ABSA-Pair	$D_{20a-14lap}$	339	0	0.00%
ABSA-Triplet	$D_{20b-14lap}$	328	0	0.00%
NER-Flat	CoNLL03	2000	0	0.00%
NER-Nested	ACE05-Ent	1050	5	0.48%
RE-RC	SemEval2010	2000	4	0.20%
RE-Triplet	CoNLL04	287	1	0.35%
EE-Trigger	ACE05-Evt	284	3	1.06%
EE-Argument	ACE05-Evt	403	0	0.00%
EE-Joint	ACE05-Evt	284	34	11.97%

Table 4: The ratio of invalid responses for each IE sub-task under the zero-shot setting. “#Prompt” is the number of test prompts. “#Invalid” indicates the number of test prompts with invalid responses under the zero-shot setting. “Ratio (%)” denotes the percentage of “#Invalid” in “#Prompt”.

379 6.2. Frequency of Target Types

380 The real-world data usually exhibits a long-tailed distribution, i.e., the fre-
 381 quency of target types varies greatly, causing the models to perform much worse
 382 on uncommon/tail types than on common/head ones. Here target types include
 383 entity types, relation types, event types, etc. In this subsection, we investigate
 384 the impact of the “frequency of target types” on GPT-4’s performance on all IE
 385 sub-tasks. We select one dataset for each sub-task with the phenomenon of fre-

quency differences and report the result of zero-shot prompts on the head types and tail types. For each dataset, We use a threshold K to determine its head and tail types. The head types are those with more than K training instances in the training set, while tail types are those with less than K training instances. Take the entity type “Person” as an example, if the number of “Person” entities in the training set is more than the threshold K , then “Person” is a head type, and vice versa, it is a tail type. The values of K corresponding to each dataset used in this section are listed in Table 5. No datasets related to ABSA are shown here since ABSA tasks involve no types.

Dataset	K	#Head	#Tail
CoNLL03	8000	2	2
ACE05-Ent	1000	3	4
SemEval2010	600	4	6
NYT-multi	500	11	13
Commodity News EE	100	7	12
CASIE	1200	2	3

Table 5: The threshold (**K**), number of head types (#Head), and number of tail types (#Tail) on different datasets.

Task	Dataset	Head	Tail	Ratio (%)
NER-Flat	CoNLL03	83.94	55.12	65.67%
NER-Nested	ACE05-Ent	27.58	12.01	43.54%
RE-RC	SemEval2010	51.28	45.18	88.09%
RE-Triplet	NYT-multi	14.50	2.30	15.89%
EE-Trigger	Commodity News EE	53.01	37.62	70.97%
EE-Joint	CASIE	15.73	17.61	111.89%

Table 6: Performance comparison of GPT-4 between head and tail types for each IE sub-task. Note that the sub-tasks not listed have no long-tail distribution. “Ratio (%)” indicates the percentages of tail types’ results with respect to head types’ results.

The results are shown in Table 6. It can be seen that the performance of

tail types is generally significantly worse than that of head types, with the only exception of EE-Joint where their performances are close. For RE-Triplet, the performance of tail types is even lower than 16% of head types’ performance. We can conclude that **GPT-4 also suffers from the long-tail problem.**

6.3. Subject-object Orders

In this subsection, we explore whether GPT-4 can distinguish the order of two entities in the RE-RC sub-task, i.e., which entity is the subject and which entity is the object. Since most relation types are not symmetric, the order of two entities is very critical. For example, the sentence “*Steven Paul Jobs was born in San Francisco on February 24, 1955.*” expresses the relational triplet $\langle \textit{Steven Paul Jobs}, \textit{born_in}, \textit{San Francisco} \rangle$, not the triplet $\langle \textit{San Francisco}, \textit{born_in}, \textit{Steven Paul Jobs} \rangle$, where the subject entity should take the first place while the object entity takes the last. For each instance of the asymmetric relation types, we swap the order of entities and check the change in prediction results. After exchanging the order, the prediction result should be changed to an empty string, which indicates no relation.

Table 7 shows zero-shot prompted results for this subsection. Our experiments found that GPT-4 has made large progress in identifying subject-object order over GPT-3.5, so we list the performances of both GPT-3.5 and GPT-4 in this scenario for comparison. It can be seen that only a small amount of GPT-3.5’s predictions (less than 30%) can tell the difference between subject and object entities, which is denoted by the ratio of swapped predictions changed to no relation. For GPT-4, however, this ratio becomes much higher on all three datasets, with up to 84.77% on the CoNLL04 dataset, yet still not satisfying enough on the other two datasets. Therefore, it can be concluded that **for the RE-RC sub-task, GPT-4 is more sensitive to the order of entities compared to GPT-3.5, but still needs further improvement to understand their subject-object relationships accurately.**

Dataset	Changed (%)		
	GPT-3.5	GPT-4	Δ
CoNLL04	28.11%	84.77%	+56.66%
NYT-multi	12.56%	46.03%	+33.47%
SemEval2010	18.02%	49.05%	+31.03%

Table 7: Statistics of changes in predicted results after swapping the order of entities. “Changed (%)” denotes the percentage of predictions changed to empty strings when swapping subject entity and object entity. Δ means the increase of GPT-4 compared to GPT-3.5.

7. Analysis of Error Types

In this section, we analyze GPT-4’s errors on all IE sub-tasks. Here we use “span” to denote the target information to be extracted, and “types” to indicate the types of target information such as entity types, relation types, event types, sentiment polarity, etc. Through manual checking, we find that the errors mainly include:

1. **Missing spans:** Missing one or more annotated target spans.
2. **Missing types:** The annotated target span is correctly answered, but its corresponding type isn’t.
3. **Unmentioned spans:** Answering the spans that do not exist within the given input text.
4. **Unannotated spans:** Answering the spans that are not annotated in the test set.
5. **Incorrect span offsets:** The offsets of the answered spans are incorrect.
6. **Undefined types:** Answering the types beyond the pre-defined types when the corresponding span is correct.
7. **Incorrect types:** The answered span is correct, while the corresponding type comes from the set of pre-defined types, but does not match the annotated type.

Here, the first two error types form the false negative samples, and error types 3 to 7 form the false positive ones. Since these error types are suitable for

Error Type	ABSA-AESC		NER-Flat		RE-Triplet		EE-Trigger	
	Num.	Ratio (%)	Num.	Ratio (%)	Num.	Ratio (%)	Num.	Ratio (%)
Missing spans	112	16.45%	443	22.74%	294	53.07%	215	46.74%
Missing types	102	14.98%	546	28.03%	27	4.87%	24	5.22%
Unmentioned spans	11	1.62%	3	0.15%	24	4.33%	0	0.00%
Unannotated spans	330	48.46%	227	11.65%	188	33.94%	178	38.70%
Incorrect span offsets	80	11.75%	204	10.47%	0	0.00%	23	5.00%
Undefined types	2	0.29%	44	2.26%	2	0.36%	0	0.00%
Incorrect types	44	6.46%	481	24.69%	19	3.43%	20	4.35%
Total	681	100.00%	1948	100.00%	554	100.00%	460	100.00%

Table 8: Statistical analysis of various error types for ABSA-AESC, NER-Flat, RE-Triplet, and EE-Trigger sub-tasks on D_{20a} -14lap, CoNLL03, CoNLL04, and ACE05-Evt datasets respectively. “Num.” indicates the occurrence number of corresponding error type, while “Ratio (%)” denotes the corresponding percentage.

all sub-tasks in Section 3.1, for convenience, we take ABSA-AESC, NER-Flat, RE-Triplet, and EE-Trigger as examples (one for each task) and statistically analyze each error type under the zero-shot setting. The ABSA tasks involve no types, so we treat the sentiment polarities of the aspects as types when analyzing ABSA-AESC. The results are shown in Table 8 and Figure 2.




Figure 2: Percentage of error types for ABSA-AESC, NER-Flat, RE-Triplet, and EE-Trigger sub-tasks on D_{20a} -14lap, CoNLL03, CoNLL04, and ACE05-Evt datasets respectively.

450 It can be seen that “Missing spans” and “Unannotated spans” are the
451 two main types of errors, accounting for more than 60% in most cases.
452 In particular, “Unannotated spans” accounts for more than 1/3 of all errors
453 in ABSA-AESC, RE-Triplet, and EE-Trigger sub-tasks, which raises concerns
454 about the quality of annotated data.

455 8. Human-like Improvement Methods

456 In the previous sections, we conduct a comprehensive performance evaluation
457 and robustness analysis on the information extraction ability of GPT-4.
458 The results suggest a significant gap between the performance of GPT-4 and the
459 SOTA methods of IE tasks. Although GPT-4’s performance has improved significantly
460 compared to GPT-3.5, the above issues, especially its performance on
461 difficult sub-tasks such as EE-Joint, restrict its application value in IE tasks. In
462 this section, we propose three novel prompt-based methods for improving the IE
463 performance of human-like LLMs like GPT-4 in the same way as to make human
464 laymen more proficient: **Task-related Knowledge Informing**, **Methodology Specifying**, and **Sufficient Extraction Reminder**. The corresponding
465 modified prompts are listed in Appendix B for more references.
466

467 8.1. Task-related Knowledge Informing

468 Compared with ordinary people, LLMs like GPT-3.5 and GPT-4 have the
469 advantage of learning extensive knowledge from pre-training corpora (covering
470 knowledge resources such as Wikipedia) and therefore have the potential to be
471 used as knowledge bases [89]. However, given the professional nature of information
472 extraction, completing these tasks may require some task-related knowledge
473 that is rarely seen in the corpus or easily confused with other knowledge, such
474 as the meaning of words like “trigger” and “argument” that have a large semantic
475 difference in the context of information extraction and other times. For a
476 layman participating in information extraction for the first time, no matter if it
477 is a human or an unfine-tuned LLM, informing it of the task-related knowledge

Task	Dataset	Base	No TKI.		MS.		SER.	
			Result	Δ	Result	Δ	Result	Δ
ABSA-AE	$D_{17-14lap}$	46.06	47.01	0.95	-	-	47.38	1.32
ABSA-OE	$D_{17-14lap}$	51.44	45.59	-5.85	-	-	50.2	-1.24
ABSA-ALSC	$D_{17-14lap}$	79.82	78.75	-1.07	-	-	78.59	-1.22
ABSA-AOE	$D_{19-14lap}$	56.57	52.87	-3.7	-	-	57.69	1.11
ABSA-AESC	$D_{20a-14lap}$	44.77	44.63	-0.14	-	-	44.73	-0.04
ABSA-Pair	$D_{20a-14lap}$	38.2	41.16	2.96	-	-	39.46	1.26
ABSA-Triplet	$D_{20b-14lap}$	35.49	33.47	-2.02	-	-	35.4	-0.09
NER-Flat	CoNLL03	72.3	74.38	2.09	67.25	-5.05	69.23	-3.06
NER-Nested	ACE05-Ent	24.68	29.96	5.28	28.98	4.3	25.81	1.13
RE-RC	SemEval2010	44.80	-	-	43.87	-0.93	44.44	-0.37
RE-Triplet	CoNLL04	26.53	-	-	28.57	2.05	24.87	-1.65
EE-Trigger	ACE05-Evt	41.62	41.99	0.36	41.16	-0.46	41.12	-0.51
EE-Argument	ACE05-Evt	29.43	27.25	-2.19	28.52	-0.91	29.03	-0.4
EE-Joint	ACE05-Evt	3.7	17.07	13.37	20.52	16.82	18.88	15.18

Table 9: Statistical analysis of the effect of three performance improvement methods. “Base” denotes the result of the basic zero-shot prompt used in Section 4, “No TKI.” indicates the result without Task-related Knowledge Informing, and “MS.” and “SER.” denotes the result of Methodology Specifying and Sufficient Extraction Reminder respectively. Δ denotes the performance gains and losses caused by the methods compared to “Base”.

in the instructions should be a basis for performing the task, without which the performance may be affected. Considering this, all prompts used in Section 4 contain the task-related knowledge required for the corresponding sub-tasks, including:

- **For NER Tasks:** Instead of giving the abbreviations of entity types, provide the model with their verbose, e.g., provide “geographical political entities” instead of “GPE”.
- **For EE Tasks:** The meanings of “trigger” and “argument” in the context of EE tasks.
- **For ABSA Tasks:** The meanings of “aspect” and “opinion” in the con-

488 text of ABSA tasks.

489 To demonstrate the effect of this prompt design method, Table 9 shows
490 the results of the zero-shot prompt without task-related knowledge. Providing
491 task-related knowledge does come into effect in some cases, especially in recog-
492 nizing subject-object orders. Surprisingly, however, half of the presented cases
493 witness a performance increase without task-related knowledge, especially EE-
494 joint, where the performance increases by 13.37 F1 value. We speculate that this
495 is because GPT-4’s ability to process long texts is limited even if the text length
496 is within its working range, which results in a decrease in its ability to under-
497 stand prompts after adding additional knowledge, especially when working on
498 complex tasks like EE-joint that require long prompts to express. In addition,
499 GPT-4 probably can understand the type abbreviations involved in the two
500 NER subtasks well (i.e., it may have frequent contact with them in pretraining)
501 to correctly classify spans to these abbreviations. Therefore, it can be concluded
502 that **providing task-related knowledge to LLMs like GPT-4 may not**
503 **improve their IE performance in some cases.** The LLMs’ long text pro-
504 cessing ability and understanding of relevant knowledge should be considered
505 when deciding whether to add it to prompts for performance improvement.

506 8.2. Methodology Specifying

507 In addition to informing laymen of the relevant task-related knowledge in
508 performing IE tasks, another common way to improve performance is to specify
509 for them a methodology to solve IE problems. In [63], the execution of IE tasks
510 is divided into two stages, each invoked by a prompt and responsible for part
511 of the task, and the result of the first stage paves the way for the solution of
512 the second stage. Inspired by that, the prompt of our method is also divided
513 into two stages with such characteristics, but we only need to send the prompt
514 once to let the LLM solve both stages. The prompt first requires the LLM to
515 check whether each given type of information to be extracted exists in the given
516 text. Then, the prompt demands the LLM to output all information belonging

517 to each existing type. Following [63], this method is not applied to ABSA tasks
518 either, because their target information has no type.

519 Table 9 shows the results of the improvement method in this section on the
520 zero-shot prompt. For most tasks, this method brings significant performance
521 gains, with up to 16.82 F1 value on the EE-Joint sub-task. It is worth noting
522 that this method generally improves more in difficult tasks, but may reduce
523 scores on simple tasks like the NER-Flat sub-task. To conclude, **this method**
524 **can be used to improve performance on difficult tasks with inspiring**
525 **results.**

526 8.3. Sufficient Extraction Reminder

527 It can be found in Section 7 that “missing spans” is one of the major er-
528 ror types in GPT-4’s output. Therefore, explicitly reminding the model in the
529 prompt to extract all the required information in the text it can is a fairly obvi-
530 ous human-like way to improve performance. Table 9 shows the results of this
531 method under the zero-shot setting. This method achieves a significant perfor-
532 mance gain on EE-Joint with a 15.18 F1 value increase but has a relatively small
533 effect on other sub-tasks. Therefore, it can be concluded that **this method is**
534 **effective but narrowly adaptable and may only be suitable for very**
535 **difficult tasks.**

536 9. Conclusion

537 In this paper, we assess the capabilities of GPT-4 from four perspectives
538 including Performance, Evaluation Criteria, Robustness, and Error Types. We
539 then proposed three prompt-based human-like Improvement Methods to en-
540 hance the models’ performance. The details and conclusions are as follows:

541 **Performance.** We first evaluate GPT-4’s performance on 16 datasets with
542 14 IE sub-tasks under the zero-shot, few-shot, and chain-of-thought scenarios,
543 and find a visible improvement over GPT-3.5 and a significant performance gap
544 compared to SOTA results.

545 ***Evaluation Criteria.*** We rethink the performance gap and find that the span
546 hard-matching strategy is not suitable for the evaluation of GPT-4, because
547 GPT-4 generates human-like responses. We propose a soft-matching strategy
548 for evaluation to reflect GPT-4’s performance more accurately.

549 ***Robustness.*** We analyze the robustness of ChatGPT on 14 IE sub-tasks from
550 three perspectives, including invalid output, frequency of target types, and error
551 types. We draw the following conclusions: 1) GPT-4 rarely outputs invalid re-
552 sponses; 2) Long-tail target types greatly affect GPT-4’s performance; 3) GPT-4
553 can understand the subject-object relationships in the RE-RC sub-task better
554 than GPT-3.5, but still needs further improvement.

555 ***Error Types.*** Through manual checking, we analyze the errors of GPT-4 and
556 summarize 7 types of errors, including Missing spans, Missing types, Unmen-
557 tioned spans, Unannotated spans, Incorrect span offsets, Undefined types, and
558 Incorrect types. We find that “Missing spans” and “Unannotated spans” are the
559 most dominant error types. The widespread presence of “Unannotated spans”
560 also raises concerns about the quality of previously annotated data and indicates
561 the possibility of annotating data with GPT-4.

562 ***Improvement Methods.*** We propose three methods to improve the IE perfor-
563 mance of LLMs like GPT-4, namely Task-related Knowledge Informing, Method-
564 ology Specifying, and Sufficient Extraction Reminder. Among them, Method-
565 ology Specifying achieves the most significant performance gains. Rich experi-
566 ments prove the effectiveness of these methods, while also showing the necessity
567 to consider the model capability, task difficulty, and characteristics when decid-
568 ing whether or not to adopt these methods.

569 CRediT authorship contribution statement

570 **Ridong Han:** Investigation, Methodology, Writing - Original Draft, Vi-
571 sualization, Project administration, Software. **Chaohao Yang:** Methodology,
572 Writing - Original Draft, Visualization, Software. **Tao Peng:** Supervision,

573 Funding acquisition. **Prayag Tiwari**: Validation. **Xiang Wan**: Resources.
574 **Lu Liu**: Validation, Funding acquisition. **Benyou Wang**: Supervision, Writ-
575 ing - Review & Editing.

576 Declaration of Competing Interest

577 The authors declare that they have no known competing financial interests or
578 personal relationships that could have appeared to influence the work reported
579 in this paper.

580 Acknowledgment

581 This work is supported by the National Natural Science Foundation of China
582 under grant No.61872163 and 61806084, and Jilin Province Key Scientific and
583 Technological Research and Development Project under grant No.20210201131GX.

584 References

- 585 [1] T. B. Brown, B. Mann, N. Ryder, M. Subbiah, J. Kaplan, P. Dhari-
586 wal, A. Neelakantan, P. Shyam, G. Sastry, A. Aspell, S. Agarwal,
587 A. Herbert-Voss, G. Krueger, T. Henighan, R. Child, A. Ramesh,
588 D. M. Ziegler, J. Wu, C. Winter, C. Hesse, M. Chen, E. Sigler,
589 M. Litwin, S. Gray, B. Chess, J. Clark, C. Berner, S. McCandlish,
590 A. Radford, I. Sutskever, D. Amodei, Language models are few-shot
591 learners, in: Advances in Neural Information Processing Systems,
592 NeurIPS, 2020. URL: [https://proceedings.neurips.cc/paper/2020/
593 hash/1457c0d6bfc4967418bfb8ac142f64a-Abstract.html](https://proceedings.neurips.cc/paper/2020/hash/1457c0d6bfc4967418bfb8ac142f64a-Abstract.html).
- 594 [2] R. Thoppilan, D. D. Freitas, J. Hall, N. Shazeer, A. Kulshreshtha,
595 H. Cheng, A. Jin, T. Bos, L. Baker, Y. Du, Y. Li, H. Lee, H. S. Zheng,
596 A. Ghafouri, M. Menegali, Y. Huang, M. Krikun, D. Lepikhin, J. Qin,
597 D. Chen, Y. Xu, Z. Chen, A. Roberts, M. Bosma, Y. Zhou, C. Chang,
598 I. Krivokon, W. Rusch, M. Pickett, K. S. Meier-Hellstern, M. R. Morris,

- 599 T. Doshi, R. D. Santos, T. Duke, J. Soraker, B. Zevenbergen, V. Prab-
600 hakaran, M. Diaz, B. Hutchinson, K. Olson, A. Molina, E. Hoffman-John,
601 J. Lee, L. Aroyo, R. Rajakumar, A. Butryna, M. Lamm, V. Kuzmina,
602 J. Fenton, A. Cohen, R. Bernstein, R. Kurzweil, B. Aguera-Arcas, C. Cui,
603 M. Croak, E. H. Chi, Q. Le, Lamda: Language models for dialog applica-
604 tions, CoRR abs/2201.08239 (2022). URL: [https://doi.org/10.48550/](https://doi.org/10.48550/arXiv.2201.08239)
605 [arXiv.2201.08239](https://doi.org/10.48550/arXiv.2201.08239). doi:10.48550/arXiv.2201.08239.
- 606 [3] A. Chowdhery, S. Narang, J. Devlin, M. Bosma, G. Mishra, A. Roberts,
607 P. Barham, H. W. Chung, C. Sutton, S. Gehrmann, P. Schuh, K. Shi,
608 S. Tsvyashchenko, J. Maynez, A. Rao, P. Barnes, Y. Tay, N. Shazeer,
609 V. Prabhakaran, E. Reif, N. Du, B. Hutchinson, R. Pope, J. Bradbury,
610 J. Austin, M. Isard, G. Gur-Ari, P. Yin, T. Duke, A. Levskaya, S. Ghe-
611 mawat, S. Dev, H. Michalewski, X. Garcia, V. Misra, K. Robinson, L. Fe-
612 dus, D. Zhou, D. Ippolito, D. Luan, H. Lim, B. Zoph, A. Spiridonov, R. Sep-
613 assi, D. Dohan, S. Agrawal, M. Omernick, A. M. Dai, T. S. Pillai, M. Pel-
614 lat, A. Lewkowycz, E. Moreira, R. Child, O. Polozov, K. Lee, Z. Zhou,
615 X. Wang, B. Saeta, M. Diaz, O. Firat, M. Catasta, J. Wei, K. Meier-
616 Hellstern, D. Eck, J. Dean, S. Petrov, N. Fiedel, Palm: Scaling language
617 modeling with pathways, Journal of Machine Learning Research 24 (2023)
618 240:1–240:113. URL: <http://jmlr.org/papers/v24/22-1144.html>.
- 619 [4] J. Wei, X. Wang, D. Schuurmans, M. Bosma, B. Ichter, F. Xia,
620 E. H. Chi, Q. V. Le, D. Zhou, Chain-of-thought prompt-
621 ing elicits reasoning in large language models, in: Advances
622 in Neural Information Processing Systems, NeurIPS, 2022.
623 URL: [http://papers.nips.cc/paper_files/paper/2022/hash/](http://papers.nips.cc/paper_files/paper/2022/hash/9d5609613524ecf4f15af0f7b31abca4-Abstract-Conference.html)
624 [9d5609613524ecf4f15af0f7b31abca4-Abstract-Conference.html](http://papers.nips.cc/paper_files/paper/2022/hash/9d5609613524ecf4f15af0f7b31abca4-Abstract-Conference.html).
- 625 [5] V. Sanh, A. Webson, C. Raffel, S. H. Bach, L. Sutawika, Z. Alyafeai,
626 A. Chaffin, A. Stiegler, A. Raja, M. Dey, M. S. Bari, C. Xu, U. Thakker,
627 S. S. Sharma, E. Szczechla, T. Kim, G. Chhablani, N. V. Nayak, D. Datta,
628 J. Chang, M. T. Jiang, H. Wang, M. Manica, S. Shen, Z. X. Yong,

- 629 H. Pandey, R. Bawden, T. Wang, T. Neeraj, J. Rozen, A. Sharma,
630 A. Santilli, T. Févry, J. A. Fries, R. Teehan, T. L. Scao, S. Bider-
631 man, L. Gao, T. Wolf, A. M. Rush, Multitask prompted training en-
632 ables zero-shot task generalization, in: Proceedings of the Tenth In-
633 ternational Conference on Learning Representations, ICLR, 2022. URL:
634 <https://openreview.net/forum?id=9Vrb9D0WI4>.
- 635 [6] L. Ouyang, J. Wu, X. Jiang, D. Almeida, C. L. Wainwright, P. Mishkin,
636 C. Zhang, S. Agarwal, K. Slama, A. Ray, J. Schulman, J. Hilton, F. Kel-
637 ton, L. Miller, M. Simens, A. Askell, P. Welinder, P. F. Christiano, J. Leike,
638 R. Lowe, Training language models to follow instructions with human feed-
639 back, in: Advances in Neural Information Processing Systems, NeurIPS,
640 2022. URL: [http://papers.nips.cc/paper_files/paper/2022/hash/](http://papers.nips.cc/paper_files/paper/2022/hash/b1efde53be364a73914f58805a001731-Abstract-Conference.html)
641 [b1efde53be364a73914f58805a001731-Abstract-Conference.html](http://papers.nips.cc/paper_files/paper/2022/hash/b1efde53be364a73914f58805a001731-Abstract-Conference.html).
- 642 [7] OpenAI, GPT-4 technical report, CoRR abs/2303.08774 (2023).
643 URL: <https://doi.org/10.48550/arXiv.2303.08774>. doi:10.48550/
644 [arXiv.2303.08774](https://doi.org/10.48550/arXiv.2303.08774).
- 645 [8] OpenAI, Introducing chatgpt, OpenAI Blog (2023). URL: [https://](https://openai.com/blog/chatgpt)
646 openai.com/blog/chatgpt.
- 647 [9] W. Jiao, W. Wang, J. Huang, X. Wang, Z. Tu, Is chatgpt
648 A good translator? A preliminary study, CoRR abs/2301.08745
649 (2023). URL: <https://doi.org/10.48550/arXiv.2301.08745>. doi:10.
650 [48550/arXiv.2301.08745](https://doi.org/10.48550/arXiv.2301.08745).
- 651 [10] Y. Bang, S. Cahyawijaya, N. Lee, W. Dai, D. Su, B. Wilie, H. Lovenia, Z. Ji,
652 T. Yu, W. Chung, Q. V. Do, Y. Xu, P. Fung, A multitask, multilingual,
653 multimodal evaluation of chatgpt on reasoning, hallucination, and interac-
654 tivity, in: Proceedings of the 13th International Joint Conference on Natu-
655 ral Language Processing and the 3rd Conference of the Asia-Pacific Chap-
656 ter of the Association for Computational Linguistics, IJCNLP, 2023, pp.

- 657 675–718. URL: <https://doi.org/10.18653/v1/2023.ijcnlp-main.45>.
658 doi:10.18653/V1/2023.IJCNLP-MAIN.45.
- 659 [11] Z. Fei, Q. Zhang, T. Gui, D. Liang, S. Wang, W. Wu, X. Huang, CQG:
660 A simple and effective controlled generation framework for multi-hop ques-
661 tion generation, in: Proceedings of the 60th Annual Meeting of the Asso-
662 ciation for Computational Linguistics, ACL, 2022, pp. 6896–6906. URL:
663 <https://doi.org/10.18653/v1/2022.acl-long.475>. doi:10.18653/v1/
664 2022.acl-long.475.
- 665 [12] S. Cao, J. Shi, Z. Yao, X. Lv, J. Yu, L. Hou, J. Li, Z. Liu, J. Xiao, Pro-
666 gram transfer for answering complex questions over knowledge bases, in:
667 Proceedings of the 60th Annual Meeting of the Association for Computa-
668 tional Linguistics, ACL, 2022, pp. 8128–8140. URL: <https://doi.org/10.18653/v1/2022.acl-long.559>. doi:10.18653/v1/2022.acl-long.559.
- 670 [13] L. Wang, W. Zhao, Z. Wei, J. Liu, Simkgc: Simple contrastive knowledge
671 graph completion with pre-trained language models, in: Proceedings of
672 the 60th Annual Meeting of the Association for Computational Linguistics,
673 ACL, 2022, pp. 4281–4294. URL: <https://doi.org/10.18653/v1/2022.acl-long.295>. doi:10.18653/v1/2022.acl-long.295.
- 675 [14] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N.
676 Gomez, L. Kaiser, I. Polosukhin, Attention is all you need, in: Ad-
677 vances in Neural Information Processing Systems, NeurIPS, 2017, pp.
678 5998–6008. URL: [https://proceedings.neurips.cc/paper/2017/hash/
679 3f5ee243547dee91fbd053c1c4a845aa-Abstract.html](https://proceedings.neurips.cc/paper/2017/hash/3f5ee243547dee91fbd053c1c4a845aa-Abstract.html).
- 680 [15] J. Devlin, M. Chang, K. Lee, K. Toutanova, BERT: pre-training of deep
681 bidirectional transformers for language understanding, in: Proceedings of
682 the 2019 Conference of the North American Chapter of the Association for
683 Computational Linguistics: Human Language Technologies, NAACL-HLT,
684 2019, pp. 4171–4186. URL: <https://doi.org/10.18653/v1/n19-1423>.
685 doi:10.18653/v1/n19-1423.

- [16] M. Lewis, Y. Liu, N. Goyal, M. Ghazvininejad, A. Mohamed, O. Levy, V. Stoyanov, L. Zettlemoyer, BART: denoising sequence-to-sequence pre-training for natural language generation, translation, and comprehension, in: Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics, ACL, 2020, pp. 7871–7880. URL: <https://doi.org/10.18653/v1/2020.acl-main.703>. doi:10.18653/v1/2020.acl-main.703.
- [17] J. Kaplan, S. McCandlish, T. Henighan, T. B. Brown, B. Chess, R. Child, S. Gray, A. Radford, J. Wu, D. Amodei, Scaling laws for neural language models, CoRR abs/2001.08361 (2020). URL: <https://doi.org/10.48550/arXiv.2001.08361>. doi:10.48550/arXiv.2001.08361.
- [18] J. Hoffmann, S. Borgeaud, A. Mensch, E. Buchatskaya, T. Cai, E. Rutherford, D. de Las Casas, L. A. Hendricks, J. Welbl, A. Clark, T. Hennigan, E. Noland, K. Millican, G. van den Driessche, B. Damoc, A. Guy, S. Osindero, K. Simonyan, E. Elsen, J. W. Rae, O. Vinyals, L. Sifre, Training compute-optimal large language models, CoRR abs/2203.15556 (2022). URL: <https://doi.org/10.48550/arXiv.2203.15556>. doi:10.48550/arXiv.2203.15556.
- [19] S. Smith, M. Patwary, B. Norick, P. LeGresley, S. Rajbhandari, J. Casper, Z. Liu, S. Prabhumoye, G. Zerveas, V. Korthikanti, E. Zheng, R. Child, R. Y. Aminabadi, J. Bernauer, X. Song, M. Shoenybi, Y. He, M. Houston, S. Tiwary, B. Catanzaro, Using deepspeed and megatron to train megatron-turing NLG 530b, A large-scale generative language model, CoRR abs/2201.11990 (2022). URL: <https://doi.org/10.48550/arXiv.2201.11990>. doi:10.48550/arXiv.2201.11990.
- [20] J. Wei, Y. Tay, R. Bommasani, C. Raffel, B. Zoph, S. Borgeaud, D. Yogatama, M. Bosma, D. Zhou, D. Metzler, E. H. Chi, T. Hashimoto, O. Vinyals, P. Liang, J. Dean, W. Fedus, Emergent abilities of large lan-

- guage models, Transactions on Machine Learning Research 2022 (2022).
URL: <https://openreview.net/forum?id=yzkSU5zdwD>.
- [21] R. Schaeffer, B. Miranda, S. Koyejo, Are emergent abilities of large language models a mirage?, in: Advances in Neural Information Processing Systems, NeurIPS, 2023.
URL: http://papers.nips.cc/paper_files/paper/2023/hash/adc98a266f45005c403b8311ca7e8bd7-Abstract-Conference.html.
- [22] J. Wei, M. Bosma, V. Y. Zhao, K. Guu, A. W. Yu, B. Lester, N. Du, A. M. Dai, Q. V. Le, Finetuned language models are zero-shot learners, in: Proceedings of the Tenth International Conference on Learning Representations, ICLR, 2022. URL: <https://openreview.net/forum?id=gEZrGCozdqR>.
- [23] T. Kojima, S. S. Gu, M. Reid, Y. Matsuo, Y. Iwasawa, Large language models are zero-shot reasoners, in: Advances in Neural Information Processing Systems, NeurIPS, 2022.
URL: http://papers.nips.cc/paper_files/paper/2022/hash/8bb0d291acd4acf06ef112099c16f326-Abstract-Conference.html.
- [24] Y. Wang, S. Mishra, P. Alipoormolabashi, Y. Kordi, A. Mirzaei, A. Naik, A. Ashok, A. S. Dhanasekaran, A. Arunkumar, D. Stap, E. Pathak, G. Karamanolakis, H. G. Lai, I. Purohit, I. Mondal, J. Anderson, K. Kuznia, K. Doshi, K. K. Pal, M. Patel, M. Moradshahi, M. Parmar, M. Purohit, N. Varshney, P. R. Kaza, P. Verma, R. S. Puri, R. Karia, S. Doshi, S. K. Sampat, S. Mishra, S. R. A, S. Patro, T. Dixit, X. Shen, Super-naturalinstructions: Generalization via declarative instructions on 1600+ NLP tasks, in: Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing, EMNLP, 2022, pp. 5085–5109. URL: <https://doi.org/10.18653/v1/2022.emnlp-main.340>. doi:10.18653/V1/2022.EMNLP-MAIN.340.

- [25] A. Radford, K. Narasimhan, T. Salimans, I. Sutskever, et al., Improving language understanding by generative pre-training (2018).
- [26] A. Radford, J. Wu, R. Child, D. Luan, D. Amodei, I. Sutskever, et al., Language models are unsupervised multitask learners, OpenAI blog 1 (2019) 9.
- [27] P. F. Christiano, J. Leike, T. B. Brown, M. Martic, S. Legg, D. Amodei, Deep reinforcement learning from human preferences, in: Advances in Neural Information Processing Systems, NeurIPS, 2017, pp. 4299–4307. URL: <https://proceedings.neurips.cc/paper/2017/hash/d5e2c0adad503c91f91df240d0cd4e49-Abstract.html>.
- [28] S. Altman, Planning for agi and beyond, OpenAI Blog (2023).
- [29] S. Bubeck, V. Chandrasekaran, R. Eldan, J. Gehrke, E. Horvitz, E. Kamar, P. Lee, Y. T. Lee, Y. Li, S. M. Lundberg, H. Nori, H. Palangi, M. T. Ribeiro, Y. Zhang, Sparks of artificial general intelligence: Early experiments with GPT-4, CoRR abs/2303.12712 (2023). URL: <https://doi.org/10.48550/arXiv.2303.12712>. doi:10.48550/arXiv.2303.12712.
- [30] X. Li, J. Feng, Y. Meng, Q. Han, F. Wu, J. Li, A unified MRC framework for named entity recognition, in: Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics, ACL, 2020, pp. 5849–5859. URL: <https://doi.org/10.18653/v1/2020.acl-main.519>. doi:10.18653/v1/2020.acl-main.519.
- [31] X. Wang, Y. Jiang, N. Bach, T. Wang, Z. Huang, F. Huang, K. Tu, Automated concatenation of embeddings for structured prediction, in: Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing, ACL-IJCNLP, 2021, pp. 2643–2660. URL: <https://doi.org/10.18653/v1/2021.acl-long.206>. doi:10.18653/v1/2021.acl-long.206.

- 770 [32] N. Ding, G. Xu, Y. Chen, X. Wang, X. Han, P. Xie, H. Zheng, Z. Liu, Few-
771 nerd: A few-shot named entity recognition dataset, in: Proceedings of the
772 59th Annual Meeting of the Association for Computational Linguistics and
773 the 11th International Joint Conference on Natural Language Processing,
774 ACL-IJCNLP, 2021, pp. 3198–3213. URL: <https://doi.org/10.18653/v1/2021.acl-long.248>. doi:10.18653/V1/2021.ACL-LONG.248.
- 776 [33] Y. Yang, X. Hu, F. Ma, S. Li, A. Liu, L. Wen, P. S. Yu, Gaussian
777 prior reinforcement learning for nested named entity recognition, in:
778 Proceedings of the IEEE International Conference on Acoustics, Speech
779 and Signal Processing, ICASSP, 2023, pp. 1–5. URL: <https://doi.org/10.1109/ICASSP49357.2023.10097163>. doi:10.1109/ICASSP49357.
780 2023.10097163.
- 782 [34] G. Nan, Z. Guo, I. Sekulic, W. Lu, Reasoning with latent structure re-
783 finement for document-level relation extraction, in: Proceedings of the
784 58th Annual Meeting of the Association for Computational Linguistics,
785 ACL, 2020, pp. 1546–1557. URL: <https://doi.org/10.18653/v1/2020.acl-main.141>. doi:10.18653/v1/2020.acl-main.141.
- 787 [35] K. Zhao, H. Xu, Y. Cheng, X. Li, K. Gao, Representation itera-
788 tive fusion based on heterogeneous graph neural network for joint entity
789 and relation extraction, Knowledge-Based Systems 219 (2021) 106888.
790 URL: <https://doi.org/10.1016/j.knosys.2021.106888>. doi:10.1016/
791 j.knosys.2021.106888.
- 792 [36] R. Han, T. Peng, J. Han, H. Cui, L. Liu, Distantly supervised relation ex-
793 traction via recursive hierarchy-interactive attention and entity-order per-
794 ception, Neural Networks 152 (2022) 191–200. URL: <https://doi.org/10.1016/j.neunet.2022.04.019>. doi:10.1016/j.neunet.2022.04.019.
- 796 [37] Y. Zhan, Z. Li, X. Zhao, C. Zhang, T. Chen, A simple overlapping re-
797 lation extraction method based on dropout, in: Proceedings of the 2022

- 798 International Joint Conference on Neural Networks, IJCNN, 2022, pp. 1–
799 8. URL: <https://doi.org/10.1109/IJCNN55064.2022.9892310>. doi:10.
800 1109/IJCNN55064.2022.9892310.
- 801 [38] B. Li, W. Ye, J. Zhang, S. Zhang, Reviewing labels: Label graph network
802 with top-k prediction set for relation extraction, in: Proceedings of the
803 Thirty-Seventh AAAI Conference on Artificial Intelligence, AAAI, 2023,
804 pp. 13051–13058. URL: <https://doi.org/10.1609/aaai.v37i11.26533>.
805 doi:10.1609/AAAI.V37I11.26533.
- 806 [39] T. Peng, R. Han, H. Cui, L. Yue, J. Han, L. Liu, Distantly super-
807 vised relation extraction using global hierarchy embeddings and local
808 probability constraints, Knowledge-Based Systems 235 (2022) 107637.
809 URL: <https://doi.org/10.1016/j.knosys.2021.107637>. doi:10.1016/
810 j.knosys.2021.107637.
- 811 [40] Z. Wang, H. Nie, W. Zheng, Y. Wang, X. Li, A novel tensor learning model
812 for joint relational triplet extraction, IEEE Transactions on Cybernetics 54
813 (2024) 2483–2494. URL: <https://doi.org/10.1109/TCYB.2023.3265851>.
814 doi:10.1109/TCYB.2023.3265851.
- 815 [41] Y. Lin, H. Ji, F. Huang, L. Wu, A joint neural model for informa-
816 tion extraction with global features, in: Proceedings of the 58th Annual
817 Meeting of the Association for Computational Linguistics, ACL, 2020, pp.
818 7999–8009. URL: <https://doi.org/10.18653/v1/2020.acl-main.713>.
819 doi:10.18653/v1/2020.acl-main.713.
- 820 [42] M. Lee, L.-K. Soon, E.-G. Siew, Effective use of graph convolution net-
821 work and contextual sub-tree for commodity news event extraction, in:
822 Proceedings of the Third Workshop on Economics and Natural Language
823 Processing, 2021, pp. 69–81.
- 824 [43] I. Hsu, K. Huang, E. Boschee, S. Miller, P. Natarajan, K. Chang, N. Peng,
825 DEGREE: A data-efficient generation-based event extraction model, in:

- 826 Proceedings of the 2022 Conference of the North American Chapter of the
827 Association for Computational Linguistics: Human Language Technologies,
828 NAACL-HLT, 2022, pp. 1890–1908. URL: <https://doi.org/10.18653/v1/2022.naacl-main.138>. doi:10.18653/v1/2022.naacl-main.138.
829
- 830 [44] Z. Chen, T. Qian, Relation-aware collaborative learning for unified aspect-
831 based sentiment analysis, in: Proceedings of the 58th Annual Meet-
832 ing of the Association for Computational Linguistics, ACL, 2020, pp.
833 3685–3694. URL: <https://doi.org/10.18653/v1/2020.acl-main.340>.
834 doi:10.18653/v1/2020.acl-main.340.
- 835 [45] H. Yan, J. Dai, T. Ji, X. Qiu, Z. Zhang, A unified generative frame-
836 work for aspect-based sentiment analysis, in: Proceedings of the 59th
837 Annual Meeting of the Association for Computational Linguistics and
838 the 11th International Joint Conference on Natural Language Processing,
839 ACL-IJCNLP, 2021, pp. 2416–2429. URL: <https://doi.org/10.18653/v1/2021.acl-long.188>.
840 doi:10.18653/v1/2021.acl-long.188.
- 841 [46] Y. Feng, Y. Rao, Y. Tang, N. Wang, H. Liu, Target-specified se-
842 quence labeling with multi-head self-attention for target-oriented opin-
843 ion words extraction, in: Proceedings of the 2021 Conference
844 of the North American Chapter of the Association for Computa-
845 tional Linguistics: Human Language Technologies, NAACL-HLT, 2021,
846 pp. 1805–1815. URL: [https://doi.org/10.18653/v1/2021.naacl-main.](https://doi.org/10.18653/v1/2021.naacl-main.145)
847 145. doi:10.18653/v1/2021.naacl-main.145.
- 848 [47] Y. Zhang, T. Peng, R. Han, J. Han, L. Yue, L. Liu, Synchronously
849 tracking entities and relations in a syntax-aware parallel architecture for
850 aspect-opinion pair extraction, Applied Intelligence 52 (2022) 15210–
851 15225. URL: <https://doi.org/10.1007/s10489-022-03286-w>. doi:10.
852 1007/s10489-022-03286-w.
- 853 [48] Y. Zhang, Y. Yang, Y. Li, B. Liang, S. Chen, Y. Dang, M. Yang,
854 R. Xu, Boundary-driven table-filling for aspect sentiment triplet ex-

- traction, in: Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing, EMNLP, 2022, pp. 6485–6498. URL: <https://doi.org/10.18653/v1/2022.emnlp-main.435>. doi:10.18653/V1/2022.EMNLP-MAIN.435.
- [49] C. Yu, T. Wu, J. Li, X. Bai, Y. Yang, Syngen: A syntactic plug-and-play module for generative aspect-based sentiment analysis, in: Proceedings of the IEEE International Conference on Acoustics, Speech and Signal Processing, ICASSP, 2023, pp. 1–5. URL: <https://doi.org/10.1109/ICASSP49357.2023.10094591>. doi:10.1109/ICASSP49357.2023.10094591.
- [50] M. U. Haque, I. Dharmadasa, Z. T. Sworna, R. N. Rajapakse, H. Ahmad, "i think this is the most disruptive technology": Exploring sentiments of chatgpt early adopters using twitter data, CoRR abs/2212.05856 (2022). URL: <https://doi.org/10.48550/arXiv.2212.05856>. doi:10.48550/arXiv.2212.05856.
- [51] S. Krügel, A. Ostermaier, M. Uhl, The moral authority of chatgpt, CoRR abs/2301.07098 (2023). URL: <https://doi.org/10.48550/arXiv.2301.07098>. doi:10.48550/arXiv.2301.07098.
- [52] R. Tang, X. Han, X. Jiang, X. Hu, Does synthetic data generation of llms help clinical text mining?, CoRR abs/2303.04360 (2023). URL: <https://doi.org/10.48550/arXiv.2303.04360>. doi:10.48550/arXiv.2303.04360.
- [53] K. Jeblick, B. Schachtner, J. Dextl, A. Mittermeier, A. T. Stüber, J. Topalis, T. Weber, P. Wesp, B. O. Sabel, J. Ricke, M. Ingrisch, Chatgpt makes medicine easy to swallow: An exploratory case study on simplified radiology reports, CoRR abs/2212.14882 (2022). URL: <https://doi.org/10.48550/arXiv.2212.14882>. doi:10.48550/arXiv.2212.14882.
- [54] S. Chen, B. H. Kann, M. B. Foote, H. J. Aerts, G. K. Savova, R. H. Mak,

883 D. S. Bitterman, The utility of chatgpt for cancer treatment information,
884 medRxiv (2023) 2023–03.

885 [55] K. Malinka, M. Peresíni, A. Firc, O. Hujnak, F. Janus, On the ed-
886 ucational impact of chatgpt: Is artificial intelligence ready to obtain a
887 university degree?, in: Proceedings of the 2023 Conference on Innova-
888 tion and Technology in Computer Science Education, ITiCSE, 2023, pp.
889 47–53. URL: [https://doi.org/10.1145/](https://doi.org/10.1145/3587102.3588827)
890 [3587102.3588827](https://doi.org/10.1145/3587102.3588827). doi:10.1145/
3587102.3588827.

891 [56] Z. Sun, A short survey of viewing large language models in legal aspect,
892 CoRR abs/2303.09136 (2023). URL: [https://doi.org/10.48550/arXiv.](https://doi.org/10.48550/arXiv.2303.09136)
893 [2303.09136](https://doi.org/10.48550/arXiv.2303.09136). doi:10.48550/arXiv.2303.09136.

894 [57] B. Zhang, D. Ding, L. Jing, How would stance detection tech-
895 niques evolve after the launch of chatgpt?, CoRR abs/2212.14548
896 (2022). URL: <https://doi.org/10.48550/arXiv.2212.14548>. doi:10.
897 [48550/arXiv.2212.14548](https://doi.org/10.48550/arXiv.2212.14548).

898 [58] B. Guo, X. Zhang, Z. Wang, M. Jiang, J. Nie, Y. Ding, J. Yue, Y. Wu,
899 How close is chatgpt to human experts? comparison corpus, evaluation,
900 and detection, CoRR abs/2301.07597 (2023). URL: [https://doi.org/](https://doi.org/10.48550/arXiv.2301.07597)
901 [10.48550/arXiv.2301.07597](https://doi.org/10.48550/arXiv.2301.07597). doi:10.48550/arXiv.2301.07597.

902 [59] T. Susnjak, Applying BERT and chatgpt for sentiment analy-
903 sis of lyme disease in scientific literature, CoRR abs/2302.06474
904 (2023). URL: <https://doi.org/10.48550/arXiv.2302.06474>. doi:10.
905 [48550/arXiv.2302.06474](https://doi.org/10.48550/arXiv.2302.06474).

906 [60] C. Qin, A. Zhang, Z. Zhang, J. Chen, M. Yasunaga, D. Yang, Is chatgpt
907 a general-purpose natural language processing task solver?, in: Proceed-
908 ings of the 2023 Conference on Empirical Methods in Natural Language
909 Processing, EMNLP, 2023, pp. 1339–1384. URL: [https://doi.org/10.](https://doi.org/10.18653/v1/2023.emnlp-main.85)
910 [18653/v1/2023.emnlp-main.85](https://doi.org/10.18653/v1/2023.emnlp-main.85). doi:10.18653/V1/2023.EMNLP-MAIN.85.

- 911 [61] Q. Zhong, L. Ding, J. Liu, B. Du, D. Tao, Can chatgpt understand
912 too? A comparative study on chatgpt and fine-tuned BERT, CoRR
913 abs/2302.10198 (2023). URL: [https://doi.org/10.48550/arXiv.2302.](https://doi.org/10.48550/arXiv.2302.10198)
914 10198. doi:10.48550/arXiv.2302.10198.
- 915 [62] N. Bian, X. Han, L. Sun, H. Lin, Y. Lu, B. He, S. Jiang, B. Dong, Chat-
916 gpt is a knowledgeable but inexperienced solver: An investigation of com-
917 monsense problem in large language models, in: Proceedings of the 2024
918 Joint International Conference on Computational Linguistics, Language
919 Resources and Evaluation, LREC/COLING, 2024, pp. 3098–3110. URL:
920 <https://aclanthology.org/2024.lrec-main.276>.
- 921 [63] X. Wei, X. Cui, N. Cheng, X. Wang, X. Zhang, S. Huang, P. Xie, J. Xu,
922 Y. Chen, M. Zhang, Y. Jiang, W. Han, Zero-shot information extraction via
923 chatting with chatgpt, CoRR abs/2302.10205 (2023). URL: [https://doi.](https://doi.org/10.48550/arXiv.2302.10205)
924 [org/10.48550/arXiv.2302.10205](https://doi.org/10.48550/arXiv.2302.10205). doi:10.48550/arXiv.2302.10205.
- 925 [64] B. Li, G. Fang, Y. Yang, Q. Wang, W. Ye, W. Zhao, S. Zhang, Evaluat-
926 ing chatgpt’s information extraction capabilities: An assessment of perfor-
927 mance, explainability, calibration, and faithfulness, CoRR abs/2304.11633
928 (2023). URL: <https://doi.org/10.48550/arXiv.2304.11633>. doi:10.
929 48550/arXiv.2304.11633.
- 930 [65] E. F. T. K. Sang, F. D. Meulder, Introduction to the conll-2003 shared
931 task: Language-independent named entity recognition, in: Proceedings of
932 the Seventh Conference on Natural Language Learning, CoNLL, 2003, pp.
933 142–147. URL: <https://aclanthology.org/W03-0419/>.
- 934 [66] G. R. Doddington, A. Mitchell, M. A. Przybocki, L. A. Ramshaw, S. M.
935 Strassel, R. M. Weischedel, The automatic content extraction (ACE) pro-
936 gram - tasks, data, and evaluation, in: Proceedings of the Fourth Interna-
937 tional Conference on Language Resources and Evaluation, LREC, 2004.
938 URL: [http://www.lrec-conf.org/proceedings/lrec2004/summaries/](http://www.lrec-conf.org/proceedings/lrec2004/summaries/5.htm)
939 5.htm.

- [67] C. Walker, S. Strassel, J. Medero, K. Maeda, Ace 2005 multilingual training corpus, Linguistic Data Consortium, Philadelphia 57 (2006) 45.
- [68] T. Ohta, Y. Tateisi, J.-D. Kim, H. Mima, J. Tsujii, The GENIA corpus: An annotated research abstract corpus in molecular biology domain, in: Proceedings of the human language technology conference, 2002, pp. 73–77.
- [69] D. Roth, W. Yih, A linear programming formulation for global inference in natural language tasks, in: Proceedings of the Eighth Conference on Computational Natural Language Learning, CoNLL, 2004, pp. 1–8. URL: <https://aclanthology.org/W04-2401/>.
- [70] X. Zeng, D. Zeng, S. He, K. Liu, J. Zhao, Extracting relational facts by an end-to-end neural model with copy mechanism, in: Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics, ACL, 2018, pp. 506–514. URL: <https://aclanthology.org/P18-1047/>. doi:10.18653/v1/P18-1047.
- [71] I. Hendrickx, S. N. Kim, Z. Kozareva, P. Nakov, D. Ó. Séaghdha, S. Padó, M. Pennacchiotti, L. Romano, S. Szpakowicz, Semeval-2010 task 8: Multi-way classification of semantic relations between pairs of nominals, in: Proceedings of the 5th International Workshop on Semantic Evaluation, SemEval@ACL, 2010, pp. 33–38. URL: <https://aclanthology.org/S10-1006/>.
- [72] T. Satyapanich, F. Ferraro, T. Finin, CASIE: extracting cybersecurity event information from text, in: Proceedings of the Thirty-Fourth AAAI Conference on Artificial Intelligence, AAAI, 2020, pp. 8749–8757. URL: <https://doi.org/10.1609/aaai.v34i05.6401>. doi:10.1609/AAAI.V34I05.6401.
- [73] M. Lee, L.-K. Soon, E.-G. Siew, L. F. Sugianto, An annotated commodity news corpus for event extraction, CoRR abs/2105.08214 (2021). URL: <https://doi.org/10.48550/arXiv.2105.08214>. doi:10.48550/arXiv.2105.08214.

- [74] W. Wang, S. J. Pan, D. Dahlmeier, X. Xiao, Coupled multi-layer attentions for co-extraction of aspect and opinion terms, in: Proceedings of the Thirty-First AAAI Conference on Artificial Intelligence, AAAI, 2017, pp. 3316–3322. URL: <https://doi.org/10.1609/aaai.v31i1.10974>. doi:10.1609/AAAI.V31I1.10974.
- [75] Z. Fan, Z. Wu, X. Dai, S. Huang, J. Chen, Target-oriented opinion words extraction with target-fused neural sequence labeling, in: Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, NAACL-HLT, 2019, pp. 2509–2518. URL: <https://doi.org/10.18653/v1/n19-1259>. doi:10.18653/v1/n19-1259.
- [76] H. Peng, L. Xu, L. Bing, F. Huang, W. Lu, L. Si, Knowing what, how and why: A near complete solution for aspect-based sentiment analysis, in: Proceedings of the Thirty-Fourth AAAI Conference on Artificial Intelligence, AAAI, 2020, pp. 8600–8607. URL: <https://doi.org/10.1609/aaai.v34i05.6383>. doi:10.1609/AAAI.V34I05.6383.
- [77] L. Xu, H. Li, W. Lu, L. Bing, Position-aware tagging for aspect sentiment triplet extraction, in: Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing, EMNLP, 2020, pp. 2339–2349. URL: <https://doi.org/10.18653/v1/2020.emnlp-main.183>. doi:10.18653/v1/2020.emnlp-main.183.
- [78] M. Pontiki, D. Galanis, J. Pavlopoulos, H. Papageorgiou, I. Androutsopoulos, S. Manandhar, Semeval-2014 task 4: Aspect based sentiment analysis, in: Proceedings of the 8th International Workshop on Semantic Evaluation, SemEval@COLING, 2014, pp. 27–35. URL: <https://doi.org/10.3115/v1/s14-2004>. doi:10.3115/v1/s14-2004.
- [79] M. Pontiki, D. Galanis, H. Papageorgiou, S. Manandhar, I. Androutsopoulos, Semeval-2015 task 12: Aspect based sentiment analysis, in: Proceedings of the 9th International Workshop on Semantic Evaluation,

- 998 SemEval@NAACL-HLT, 2015, pp. 486–495. URL: <https://doi.org/10.18653/v1/s15-2082>. doi:10.18653/v1/s15-2082.
- 999
- 1000 [80] M. Pontiki, D. Galanis, H. Papageorgiou, I. Androutsopoulos, S. Man-
- 1001 andhar, M. Al-Smadi, M. Al-Ayyoub, Y. Zhao, B. Qin, O. D. Clercq,
- 1002 V. Hoste, M. Apidianaki, X. Tannier, N. V. Loukachevitch, E. V. Kotel-
- 1003 nikov, N. Bel, S. M. J. Zafra, G. Eryigit, Semeval-2016 task 5: Aspect based
- 1004 sentiment analysis, in: Proceedings of the 10th International Workshop
- 1005 on Semantic Evaluation, SemEval@NAACL-HLT, 2016, pp. 19–30. URL:
- 1006 <https://doi.org/10.18653/v1/s16-1002>. doi:10.18653/v1/s16-1002.
- 1007 [81] H. Lv, J. Liu, H. Wang, Y. Wang, J. Luo, Y. Liu, Efficient hybrid generation
- 1008 framework for aspect-based sentiment analysis, in: Proceedings of the 17th
- 1009 Conference of the European Chapter of the Association for Computational
- 1010 Linguistics, EACL, 2023, pp. 1007–1018. URL: <https://doi.org/10.18653/v1/2023.eacl-main.71>. doi:10.18653/v1/2023.EACL-MAIN.71.
- 1011
- 1012 [82] Y. Mao, Y. Shen, C. Yu, L. Cai, A joint training dual-mrc frame-
- 1013 work for aspect based sentiment analysis, in: Proceedings of the
- 1014 Thirty-Fifth AAAI Conference on Artificial Intelligence, AAAI, 2021,
- 1015 pp. 13543–13551. URL: <https://doi.org/10.1609/aaai.v35i15.17597>.
- 1016 doi:10.1609/AAAI.V35I15.17597.
- 1017 [83] W. Zhang, X. Li, Y. Deng, L. Bing, W. Lam, Towards generative aspect-
- 1018 based sentiment analysis, in: Proceedings of the 59th Annual Meeting of
- 1019 the Association for Computational Linguistics and the 11th International
- 1020 Joint Conference on Natural Language Processing, ACL-IJCNLP, 2021,
- 1021 pp. 504–510. URL: <https://doi.org/10.18653/v1/2021.acl-short.64>.
- 1022 doi:10.18653/v1/2021.acl-short.64.
- 1023 [84] J. Lou, Y. Lu, D. Dai, W. Jia, H. Lin, X. Han, L. Sun, H. Wu, Universal
- 1024 information extraction as unified semantic matching, in: Proceedings of the
- 1025 Thirty-Seventh AAAI Conference on Artificial Intelligence, AAAI, 2023,

- 1026 pp. 13318–13326. URL: <https://doi.org/10.1609/aaai.v37i11.26563>.
 1027 doi:10.1609/AAAI.V37I11.26563.
- 1028 [85] X. Wang, W. Zhou, C. Zu, H. Xia, T. Chen, Y. Zhang, R. Zheng,
 1029 J. Ye, Q. Zhang, T. Gui, J. Kang, J. Yang, S. Li, C. Du, Instructuie:
 1030 Multi-task instruction tuning for unified information extraction, CoRR
 1031 abs/2304.08085 (2023). URL: [https://doi.org/10.48550/arXiv.2304.](https://doi.org/10.48550/arXiv.2304.08085)
 1032 08085. doi:10.48550/arXiv.2304.08085.
- 1033 [86] C. Liu, F. Zhao, Y. Kang, J. Zhang, X. Zhou, C. Sun, K. Kuang, F. Wu,
 1034 Rexuie: A recursive method with explicit schema instructor for univer-
 1035 sal information extraction, in: Findings of the Association for Computa-
 1036 tional Linguistics: EMNLP, 2023, pp. 15342–15359. URL: [https://doi.](https://doi.org/10.18653/v1/2023.findings-emnlp.1024)
 1037 [org/10.18653/v1/2023.findings-emnlp.1024](https://doi.org/10.18653/v1/2023.findings-emnlp.1024). doi:10.18653/V1/2023.
 1038 FINDINGS-EMNLP.1024.
- 1039 [87] S. Wadhwa, S. Amir, B. C. Wallace, Revisiting relation extraction in
 1040 the era of large language models, in: Proceedings of the 61st Annual
 1041 Meeting of the Association for Computational Linguistics, ACL, 2023,
 1042 pp. 15566–15589. URL: [https://doi.org/10.18653/v1/2023.acl-long.](https://doi.org/10.18653/v1/2023.acl-long.868)
 1043 868. doi:10.18653/V1/2023.ACL-LONG.868.
- 1044 [88] Y. Lu, Q. Liu, D. Dai, X. Xiao, H. Lin, X. Han, L. Sun, H. Wu, Unified
 1045 structure generation for universal information extraction, in: Proceedings
 1046 of the 60th Annual Meeting of the Association for Computational Linguis-
 1047 tics, ACL, 2022, pp. 5755–5772. URL: [https://doi.org/10.18653/v1/](https://doi.org/10.18653/v1/2022.acl-long.395)
 1048 2022.acl-long.395. doi:10.18653/v1/2022.acl-long.395.
- 1049 [89] Y. Tan, D. Min, Y. Li, W. Li, N. Hu, Y. Chen, G. Qi, Can chatgpt replace
 1050 traditional KBQA models? an in-depth analysis of the question answering
 1051 performance of the GPT LLM family, in: Proceedings of the 22nd Interna-
 1052 tional Semantic Web Conference, ISWC, volume 14265 of *Lecture Notes in*
 1053 *Computer Science*, 2023, pp. 348–367. URL: [https://doi.org/10.1007/](https://doi.org/10.1007/978-3-031-47240-4_19)
 1054 978-3-031-47240-4_19. doi:10.1007/978-3-031-47240-4_19.

1055 **Appendix A. More Performance Results**

1056 We report more performances of GPT-4 for all 14 sub-tasks in Section 3.1.
1057 Under few-shot and chain-of-thought scenarios, we design three different prompts
1058 and show their mean and respective results. The results are shown in Table B.10.

1059 **Appendix B. Example of input prompts**

1060 We show the zero-shot prompts, few-shot ICL prompts, few-shot COT prompts,
1061 and the prompts modified by the three improvement methods of the NER-Flat
1062 sub-task on the CoNLL03 dataset. Prompts for other datasets/tasks are similar.

Task	Dataset	Zero-shot	ICL				COT			
			prompt 1	prompt 2	prompt 3	mean	prompt 1	prompt 2	prompt 3	mean
ABSA-AE	$D_{17-14lap}$	46.06	55.82	48.66	54.06	52.85	55.78	56.95	55.04	55.92
	$D_{17-14res}$	64.74	72.53	73.17	75.80	73.84	76.67	76.08	74.38	75.71
	$D_{17-15res}$	48.71	58.71	60.19	54.53	57.81	59.02	58.13	57.96	58.37
ABSA-OE	$D_{17-14lap}$	51.44	61.72	53.47	56.82	57.34	41.92	34.00	41.86	39.26
	$D_{17-14res}$	62.68	62.05	74.79	73.86	70.24	57.26	61.22	67.28	61.92
	$D_{17-15res}$	49.49	57.50	54.88	58.75	57.04	45.99	55.06	51.27	50.77
ABSA-ALSC	$D_{17-14lap}$	79.82	80.28	79.51	81.96	80.58	79.05	80.73	79.97	79.92
	$D_{17-14res}$	86.46	85.49	85.84	84.16	85.16	87.52	88.23	86.37	87.37
	$D_{17-15res}$	87.06	87.99	87.43	87.06	87.49	89.65	88.17	87.99	88.60
ABSA-AOE	$D_{19-14lap}$	56.57	55.69	60.13	70.68	62.17	55.77	56.64	56.02	56.14
	$D_{19-14res}$	67.02	74.08	67.76	73.22	71.69	59.38	62.44	65.53	62.45
	$D_{19-15res}$	60.57	63.28	65.11	66.09	64.83	54.30	56.44	56.49	55.74
	$D_{19-16res}$	65.13	70.03	74.32	70.57	71.64	60.97	60.21	59.08	60.09
ABSA-AESC	$D_{20a-14lap}$	44.77	53.44	51.74	52.92	52.70	53.92	50.14	48.59	50.88
	$D_{20a-14res}$	59.94	65.29	66.13	66.38	65.93	65.21	62.23	63.61	63.69
	$D_{20a-15res}$	61.33	64.59	65.75	66.80	65.71	64.75	64.32	65.14	64.74
	$D_{20a-16res}$	56.51	61.58	65.48	63.07	63.38	66.67	65.39	65.57	65.88
ABSA-Pair	$D_{20a-14lap}$	38.20	41.83	42.25	44.31	42.80	42.35	39.78	40.61	40.91
	$D_{20a-14res}$	48.15	58.11	60.40	60.60	59.70	52.89	51.90	61.47	55.42
	$D_{20a-15res}$	50.05	56.62	54.29	55.66	55.52	52.70	50.95	51.66	51.77
	$D_{20a-16res}$	49.91	54.38	62.08	57.33	57.93	57.79	57.51	57.25	57.52
ABSA-Triplet	$D_{20b-14lap}$	35.49	41.30	40.31	38.93	40.18	37.90	34.91	36.08	36.30
	$D_{20b-14res}$	53.20	53.58	58.03	58.37	56.66	51.44	50.30	56.23	52.66
	$D_{20b-15res}$	51.26	50.43	50.80	52.71	51.31	47.22	46.17	46.50	46.63
	$D_{20b-16res}$	54.82	58.66	57.96	56.46	57.70	54.53	56.26	52.69	54.49
NER-Flat	CoNLL03	72.30	77.10	78.91	79.48	78.50	76.83	73.93	77.81	76.19
	FewNERD	47.84	50.90	50.41	46.15	49.15	52.09	50.05	49.70	50.61
NER-Nested	ACE04	31.43	43.38	42.46	44.21	43.35	45.96	45.98	46.14	46.03
	ACE05-Ent	24.68	44.60	40.51	48.93	44.68	41.45	39.12	43.83	41.46
	GENIA	46.22	57.66	57.32	55.94	56.97	55.87	54.59	53.24	54.57
RE-RC	CoNLL04	82.07	92.31	92.53	93.93	92.92	-	-	-	-
	NYT-multi	47.79	54.33	55.84	53.24	54.47	-	-	-	-
	SemEval2010	44.80	44.26	44.61	46.07	44.98	-	-	-	-
RE-Triplet	CoNLL04	26.53	35.26	33.05	43.53	37.28	34.66	35.14	37.84	35.88
	NYT-multi	12.74	23.82	15.44	19.40	19.55	14.26	11.75	14.21	13.41
	SemEval2010	5.09	15.82	16.02	18.56	16.80	-	-	-	-
EE-Trigger	ACE05-Evt	41.62	50.75	46.61	44.54	47.30	25.03	26.93	30.24	27.40
	ACE05+	43.24	47.25	47.58	46.69	47.17	23.20	28.50	28.08	26.59
	CASIE	24.98	33.12	24.33	24.04	27.17	19.78	15.38	15.61	16.92
	Commodity News EE	47.61	47.58	59.03	56.98	54.53	20.47	33.20	29.23	27.63
EE-Argument	ACE05-Evt	29.43	33.27	33.40	36.55	34.40	-	-	-	-
	ACE05+	29.72	34.47	39.39	35.09	36.32	-	-	-	-
	CASIE	30.99	31.62	31.82	30.52	31.32	-	-	-	-
	Commodity News EE	16.57	24.94	18.48	28.59	24.00	-	-	-	-
EE-Joint	ACE05-Evt	3.70	24.37	25.37	26.81	25.52	-	-	-	-
	ACE05+	3.51	22.44	29.51	26.61	26.19	-	-	-	-
	CASIE	16.78	28.53	28.83	25.89	27.75	-	-	-	-
	Commodity News EE	2.14	18.49	17.84	25.99	20.77	-	-	-	-

Table B.10: More performance results.

NER-Flat zero-shot prompt on the CoNLL03 dataset

prompt:

Considering 4 types of named entities including “organization”, “person”, “location” and “miscellaneous”, recognize all named entities in the given sentence.

Answer in the format of '[“entity_type 1”, “entity_name 1”], [“entity_type 2”, “entity_name 2”], ...' without any explanation. If no entity exists, then just answer “[]”.

Given sentence:

“Results of semifinals in the Mahindra International squash tournament on Friday :”

Expected Output:

[“miscellaneous”, “Mahindra International”]

Example NER-Flat few-shot ICL prompt on the CoNLL03 dataset

prompt:

Considering 4 types of named entities including “organization”, “person”, “location” and “miscellaneous”, recognize all named entities in the given sentence.

Answer in the format of '[“entity_type 1”, “entity_name 1”], [“entity_type 2”, “entity_name 2”], ...' without any explanation. If no entity exists, then just answer “[]”.

Sentence:

“Rapid Wien 5 0 5 0 3 3 5”

Answer:

[“organization”, “Rapid Wien”]

Sentence:

“Iran accuses Iraq of ceasefire violations .”

Answer:

[“location”, “Iran”], [“location”, “Iraq”]

... (More examples are omitted here.)

Sentence:

“The team is as follows :”

Answer:

[]

Sentence:

“Results of semifinals in the Mahindra International squash tournament on Friday :”

Answer:

Expected Output:

[“miscellaneous”, “Mahindra International”]

Example NER-Flat few-shot COT prompt on the CoNLL03 dataset

prompt:

Considering 4 types of named entities including “organization”, “person”, “location” and “miscellaneous”, recognize all named entities in the given sentence.

Answer in the format of '[“entity_type 1”, “entity_name 1”], [“entity_type 2”, “entity_name 2”], ...'. If no entity exists, then just answer “[]”. Now let’s think step by step. Focus your answers at the end of your response and don’t print them out during your thinking. If there are multiple possible types of an entity, answer the one with the highest probability.

Sentence:

“Rapid Wien 5 0 5 0 3 3 5”

Answer:

The sentence comes from the scoreline of a soccer match. “Rapid Wien” refers to the soccer team “SK Rapid Wien”, which corresponds to the “organization” in the given entity types. So, answer: [“organization”, “Rapid Wien”]

Sentence:

“Iran accuses Iraq of ceasefire violations .”

Answer:

The “Iran” and “Iraq” are all countries. The country corresponds to the “location” in the given entity types. So, answer: [“location”, “Iran”], [“location”, “Iraq”]

... (More examples are omitted here.)

Sentence:

“The team is as follows :”

Answer:

The sentence does not involve any entity of the given entity type. So, answer: []

Your sentence:

“Results of semifinals in the Mahindra International squash tournament on Friday :”

Answer:

Expected Output:

“Mahindra International” is a sports event, which fits into “miscellaneous” since it cannot be categorized into any other given types. So, answer: [“miscellaneous”, “Mahindra International”]

NER-Flat zero-shot prompt with no Task-related Knowledge Informing on the CoNLL03 dataset

prompt:

Considering 4 types of named entities including “ORG”, “PER”, “LOC” and “MISC”, recognize all named entities in the given sentence.

Answer in the format of '[“entity_type 1”, “entity_name 1”], [“entity_type 2”, “entity_name 2”], ...' without any explanation. If no entity exists, then just answer “[]”.

Given sentence:

“Results of semifinals in the Mahindra International squash tournament on Friday :”

Expected Output:

[“miscellaneous”, “Mahindra International”]

NER-Flat zero-shot prompt with Methodology Specifying on the CoNLL03 dataset

prompt:

For each of the following entity types, consider whether there are some entities of this entity type that are present in the given sentence. If present, output all entities belonging to this type in the sentence in the format of '[“entity_type 1”, “entity_name 1”], [“entity_type 2”, “entity_name 2”], ...' without any explanation, otherwise output nothing. Your answer for each existing type should take one line in response.

Given sentence:

“Results of semifinals in the Mahindra International squash tournament on Friday :”

Given entity types: “[“organization”, “person”, “location”, “miscellaneous”]”

Expected Output:

[“miscellaneous”, “Mahindra International”]

NER-Flat zero-shot prompt with Sufficient Extraction Reminder on the CoNLL03 dataset

prompt:

Considering 4 types of named entities including “organization”, “person”, “location” and “miscellaneous”, recognize all named entities in the given sentence.

Answer in the format of '[“entity_type 1”, “entity_name 1”], [“entity_type 2”, “entity_name 2”], ...' without any explanation. If no entity exists, then just answer “[]”. Please find all possible entities for each entity type as you can. If an entity appears multiple times in the text, answer it for the same number of times.

Given sentence:

“Results of semifinals in the Mahindra International squash tournament on Friday :”

Expected Output:

[“miscellaneous”, “Mahindra International”]



[Click here to access/download](#)

**LaTeX Source Files (Word or Latex - Please include all
files (editable format only) required for production)**

Latex-GPT-Evaluation.zip



Declaration of interests

☒The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

☐The authors declare the following financial interests/personal relationships which may be considered as potential competing interests: